

Implementing Dependent Types in `pi-forall`

Stephanie Weirich

June 21, 2022

1 Overview and Goals

These lecture notes describe the design of a minimal dependently-typed language called “`pi-forall`” and walk through the implementation of its type checker. They are based on lectures given at the *Oregon Programming Languages Summer School* during June 2022 and derived from earlier lectures from summer schools during 2014 and 2013.¹

What do I expect from you, dear reader? This discussion assumes a familiarity with the basics of the lambda calculus, including its standard operations (alpha-equivalence, substitution, evaluation) and the basics of type systems (especially their specification using inference rules). For background on this material, I recommend the textbooks [Pie02, Har16].

Furthermore, these notes also refer to an implementation of a demo type checker and assume basic knowledge of the Haskell programming language. This implementation is available at <https://github.com/sweirich/pi-forall> on branch 2022. As you study these notes, I encourage you to download this repository, read through its source code, and experiment with it. Detailed installation instructions are available with the repository.

What do these notes cover? These notes are broken into several sections that incrementally build up the design and implementation of a type checker for a dependently-typed programming language.

If you are looking at the repository, you’ll see that it includes several incrementally more expressive implementations of `pi-forall` in separate subdirectories. These implementations build on each other (each is an extension of the previous version) and are summarized in Figure 1. As you read each chapter, refer to its corresponding implementation to see how the features described in that chapter can be implemented.

- Section 2 starts with some examples of the full `pi-forall` language so that you can see how the pieces fit together.

¹NOTE: if the date on this document is before June 27, 2022, then these notes are still in

Key feature	<code>pi-forall</code> version	Section
Core system	<code>version1</code>	Sections 3, 4, and 5
Equality	<code>version2</code>	Sections 6 and 7
Irrelevance	<code>version3</code>	Section 8
Datatypes	<code>full</code>	Sections 2 and 9

Figure 1: Connection between sections and `pi-forall` versions

- Section 3 presents the mathematical specification of the core `pi-forall` type system including its concrete syntax (as found in `pi-forall` source files), abstract syntax (represented with a Haskell datatype), and core typing rules (written using standard mathematical notation). This initial type system is simple and declarative. In other words, its *specifies* what terms should type check, but cannot be directly implemented.
- Section 4, reformulates the typing rules so that they are *syntax-directed* and correspond to a type checking algorithm. The key idea of this section is to recast the typing rules as a *bidirectional type system*.
- Section 5 introduces the core `pi-forall` implementation and walks through the type checker found in `version1`, the Haskell implementation of the typing rules discussed in Section 4. This section also shows how the `Unbound` library can assist with the implementation of variable binding and automatically derive operations for capture-avoiding substitution and alpha-equivalence. Finally, the section describes a monadic structure for the type checker, and how it can help with the production of error messages, the primary purpose of a type checker.
- Section 6 discusses the role of definitional equality in dependently typed languages. After motivating examples, it presents both a specification of when terms are equal and a semi-decidable algorithm that can be incorporated into the type checker.
- An important feature of dependently-typed languages is the ability for run-time tests to be reflected into the type system. Section 7 shows how to extend `pi-forall` with a simple form of dependent pattern matching. We look at this feature in two different ways: First how to reflect the gain in information about boolean values in each branch of an if expression. Second, how to reflect the equality represented via propositional equality.
- Section 8 introduces the idea of tracking the *relevance* of arguments. Relevance tracking enables parts of the program to be identified as “compile-time only” and thus erased before execution. It also identifies parts of terms that can be ignored when deciding equivalence.

draft form. If you find any errors or typos, please send pull requests to <https://github.com/sweirich/pi-forall>.

- Finally, Section 9 introduces arbitrary datatypes and generalizes dependent pattern matching. This section combines the features introduced in the previous sections (Sections 6, 7, and 8) into a single unified language feature.

What do these notes not cover? The goal of this notes is to provide an introductory overview to the implementation of (dependent) type theories and type checkers. As a result, there are many related topics that are not included here. Furthermore, because these notes come with a reference implementation, they emphasize only a single point in a rich design space.

For a broader view, section 10 includes a discussion of alternative tutorials, and describes how this one differs from other approaches. The key differences are summarized below:

- For simplicity, the `pi-forall` language does not enforce termination through type checking. Implementing a proof system like Agda or Coq requires additional structure, including universe levels and bounded recursion.
- Many implementations of dependent type theories use *normalization-by-evaluation* to decide whether types are equivalent during type checking. `pi-forall` uses an alternative approach based on weak-head-normalization, defined using substitution. This approach is closer to λ -calculus theory, but can be less efficient in practice.
- Also for simplicity, `pi-forall` does not attempt to infer missing arguments using unification or other means. As a result, example programs in `pi-forall` are significantly more verbose than in other languages.
- This implementation relies on the `Unbound` library for variable binding, alpha-equivalence and substitution. As a result, these operations can be automatically derived. A more common approach is to use de Bruijn indices.
- Recent work on cubical type theory, higher-inductive types and univalence is not covered here. Indeed, the rules for dependent pattern matching that we present will be incompatible with such extension.

What do I want you to get out of all of this?

1. An understanding of how to translate mathematical specifications of type systems and logics into code, i.e. how to represent the syntax of a programming language and how to implement a type checker. More generally, this involves techniques for turning a declarative specification of a system of judgments into an algorithm that determines whether the judgement holds.
2. Exposure to the design choices of dependently-typed languages. In this respect, the goal is breadth not depth. As a result, I will provide *simple*

solutions to some of the problems that we face and sidestep other problems entirely. Because solutions are chosen for simplicity, 10 includes pointers to related work if you want to go deeper.

3. Experience with the Haskell programming language. I think Haskell is an awesome tool for this sort of work and I want to demonstrate how its features (monads, generic programming) are well-suited for this task.
4. A tool that you can use as a basis for experimentation. When you design your language, how do you know what programs you can and cannot express? Having an implementation lets you work out (smallish) examples and will help convince you (and your reviewers) that you are developing something useful. Please use `pi-forall` as a starting point for your ideas.
5. Templates and tools for writing about type systems. The source files for these lecture notes are available in the same repository as the demo implementation and I encourage you to check them out.² Building these notes requires Ott [SNO⁺07], a tool specifically tailored for typesetting type systems and mathematical specifications of programming languages.

2 A taste of `pi-forall`

Before diving into the design of the `pi-forall` language, we'll start with a few motivating examples for dependent types. This way, you'll get to learn a bit about the syntax of `pi-forall` and anticipate the definitions that are to come.

One example that dependently-typed languages are really good at is tracking the length of lists.

In `pi-forall`, we can define a datatype for length-indexed lists (often called vectors) as follows.³

```
data Vec (A : Type) (n : Nat) : Type where
  Nil of [n = Zero]
  Cons of [m:Nat] (A) (Vec A m) [n = Succ m]
```

The type `Vec A n` has two parameters: `A` the type of elements in the list and `n`, a natural number indicating the length of the list. This datatype also has two constructors: `Nil` and `Cons`, corresponding to empty and non empty lists. The `Nil` constructor constrains the length parameter to be `Zero` when it is used and the `Cons` constructor constrains it to be one more than the length of the tail of the list. In the definition of `Cons`, the `m` parameter is the length of the tail of the list. This parameter is written in square brackets to indicate that it should be *irrelevant*: this value is for type checking only and functions should not use it.

For example, we can use `pi-forall` to check that the list below contains exactly three three boolean values.

²See the `doc` subdirectory in <https://github.com/sweirich/pi-forall>.

³In the `full` implementation, you can find these definitions in the file `Vec.pi`.

```
v3 : Vec Bool 3
v3 = Cons [2] True (Cons [1] False (Cons [0] False Nil))
```

Above, `pi-forall` never infers arguments, so we must always supply the length of the tail whenever we use `Cons`.

We can also map over length indexed lists.

```
map : [A:Type] → [B:Type] → [n:Nat] → (A → B) → Vec A n → Vec B n
map = λ[A][B][n] f v.
  case v of
    Nil → Nil
    Cons [m] x xs → Cons [m] (f x) (map[A][B][m] f xs)
```

As `pi-forall` doesn't infer arguments, the `map` function needs to take `A`, `B` and `n` explicitly at the beginning of the function definition and they need to be provided as arguments to the recursive call.

Whenever we call `map`, we also need to supply the appropriate parameters for the function and vector that we are giving to `map`. For example, to map the boolean `not` function, of type `Bool → Bool`, we need to provide these two `Bools` along with the length of the vector.

```
v4 : Vec Bool 3
v4 = map [Bool][Bool][3] not v3
```

Because we are statically tracking the length of vectors, we can write a safe indexing operation. We index into the vector using an index where the type bounds its range. Below, the `Fin` type has a parameter that states what length of vector it is appropriate for.

```
data Fin (n : Nat) : Type where
  Zero of [m:Nat][n = Succ m]
  Succ of [m:Nat](Fin m)[n = Succ m]
```

For example, the type `Fin 3` has exactly three values:

```
x1 : Fin 3
x1 = Succ [2] (Succ [1] (Zero [0]))

x2 : Fin 3
x2 = Succ [2] (Zero [1])

x3 : Fin 3
x3 = Zero [2]
```

With the `Fin` type, we can safely index vectors. The following function is guaranteed to return a result because the index is within range.

```
nth : [a:Type] → [n:Nat] → Vec a n → Fin n → a
nth = λ[a][n] v f. case f of
  Zero [m] → case v of
    Cons [m'] x xs → x
  Succ [m] f' → case v of
    Cons [m'] x xs → nth [a][m] xs f'
```

Note how, in the case analysis above, neither `case` requires a branch for the `Nil` constructor. The `pi-forall` type checker can verify that the `Cons` case is exhaustive and that the `Nil` case would be inaccessible.

You may have noticed that some of the arguments to functions and constructors in the examples above are enclosed in square brackets. These arguments are *irrelevant* in `pi-forall`. What this means is that the compiler can erase these arguments before running the program and can ignore these arguments when comparing types for equality.

2.1 Homework: Check out `pi-forall`

The files in the directory `full/pi/` demonstrate the full power of the `pi-forall` language. Take a little time to familiarize yourself with some of these modules (such as `Fin.pi` and `Vec.pi`) and compare them to similar ones that you might have seen in Agda, Haskell or elsewhere.

To run the type checker on these modules, make sure that you have first compiled the implementation using `stack build` at the terminal. Then, to run `pi-forall` on a source file, such as `Fin.pi`, you can issue the command `stack exec - pi-forall Fin.pi` in the `full/pi` subdirectory. If the file type checks, you'll see the contents of the file displayed to the terminal window. Otherwise, you'll see an error message from the type checker.

For a more extensive series of examples, start with `pi/Lambda0.pi` for an interpreter for a small lambda calculus and then compare it with the implementations in `pi/Lambda1.pi` and `pi/Lambda2.pi`. These two versions index the expression type with the scope depth (a natural number) and the expression type to eliminate runtime scope and type errors from the execution of the interpreter.

2.1.1 Homework (Church encodings)

The file `full/test/NatChurch.pi` is a start at a Church encoding of natural numbers. Replace the `TRUSTME`s in this file so that it compiles. For example, one task in this file is to define a predecessor function. By replacing the `TRUSTME` with `Ref1` below, you'll be able to force the type checker to normalize both sides of the equality. The code will only type check if you have defined `pred` correctly (on this example).

```
test_pred : pred two = one
test_pred = TRUSTME -- replace with Ref1
```

3 A Simple Core Language

Now let's turn to the design and implementation of the `pi-forall` language itself. We'll start with a small core, and then incrementally add features throughout the rest of this tutorial.

Consider this simple dependently-typed lambda calculus, shown below. What

should it contain? At the bare minimum we start with the following five forms:

a, b, A, B	$::=$	x	variables
		$\lambda x. a$	lambda expressions (anonymous functions)
		$a \ b$	function applications
		$(x : A) \rightarrow B$	dependent function type, aka Π
		Type	the ‘type’ of types

As in many dependently-typed languages, we have the *same* syntax for both expressions and types. For clarity, I’ll used lowercase letters a for expressions and uppercase letters for their types A .

Note that λ and Π above are *binding forms*. They bind the variable x in a and B respectively.

3.1 When do expressions in this language type check?

We define the type system for this language using an inductive relation shown in Figure 2. This relation is between an expression a , its type A , and a typing context Γ .

$$\boxed{\Gamma \vdash a : A}$$

A typing context Γ is an ordered list of assumptions about variables and their types.

$$\Gamma ::= \emptyset \mid \Gamma, x : A$$

We will assume that each of the variables in this list are distinct from each other, so that there will always be at most one assumption about any variable’s type.⁴

An initial set of typing rules: Variables and Functions Consider the first two rules shown in Figure 2, rule T-VAR and rule T-LAMBDA. The first rule states that the types of variables are recorded in the typing context. The premise $x : A \in \Gamma$ requires us to find an association between x and the type A in the list Γ .

The second rule, for type checking functions, introduces a new variable into the context when we type check the body of a λ -expression. It also requires A , the type of the function’s argument to be a valid type.

Example: Polymorphic identity functions Note that in rule T-LAMBDA, the parameter x is allowed to appear in the result type of the function, B . Why is this useful? Well, it gives us *parametric polymorphism* automatically. In Haskell, we write the identity function as follows, annotating it with a polymorphic type.

⁴On paper, this assumption is reasonable because we always extend the context with variables that come from binders, and we can always rename bound variables so that they differ

$$\boxed{\Gamma \vdash a : A} \quad (Core\ type\ system)$$

$$\begin{array}{c}
\text{T-VAR} \\
\frac{x : A \in \Gamma}{\Gamma \vdash x : A}
\end{array}
\quad
\begin{array}{c}
\text{T-LAMBDA} \\
\frac{\Gamma, x : A \vdash a : B}{\Gamma \vdash \lambda x. a : (x : A) \rightarrow B}
\end{array}
\quad
\begin{array}{c}
\text{T-APP} \\
\frac{\Gamma \vdash a : (x : A) \rightarrow B \quad \Gamma \vdash b : A}{\Gamma \vdash a\ b : B[b/x]}
\end{array}$$

$$\begin{array}{c}
\text{T-PI} \\
\frac{\Gamma \vdash A : \mathbf{Type} \quad \Gamma, x : A \vdash B : \mathbf{Type}}{\Gamma \vdash (x : A) \rightarrow B : \mathbf{Type}}
\end{array}
\quad
\begin{array}{c}
\text{T-TYPE} \\
\frac{}{\Gamma \vdash \mathbf{Type} : \mathbf{Type}}
\end{array}$$

Figure 2: Typing rules for core system

```

id :: forall a. a → a
id = λy → y

```

Because the type of `id` is generic we can apply this function to any type of argument.

We can also write a polymorphic identity function in `pi-forall`, as follows.

```

id : (x:Type) → (y : x) → x
id = λx. λy. y

```

This definition is similar to Haskell, but with a few modifications. First, `pi-forall` uses a single colon for type annotations and a `.` for function binding instead of Haskell’s quirky `::` and `→`. Second, the `pi-forall` definition of `id` uses two lambdas: one for the “type” argument x and one for the “term” argument y . In `pi-forall`, there is no syntactic difference between these arguments: both are arguments to the identity function.

Finally, the `pi-forall` type of `id` uses the dependent function type, with general form $(x : A) \rightarrow B$. This form is like Haskell’s usual function type $A \rightarrow B$, except that we can name the argument x so that it can be referred to in the body of the type B . In `id`, dependency allows the type argument x to be used later as the type of y . We call types of this form Π -types. (Often dependent function types are written as $\Pi x : A. B$ in formalizations of dependent type theory.⁵)

The fact that the type of x is **Type** means that x plays the role of a type variable, such as `a` in the Haskell type. Because we don’t have a syntactic distinction between types and terms, in `pi-forall` we say that “types” are

from other variables in scope. Any implementation of the type system will need to somehow make sure that this invariant is maintained. We’ll do this in Section 5 using a freshness monad from the `Unbound` library.

⁵The terminology for these types is muddled: sometimes they are called dependent function types and sometimes they are called dependent product types. We use the non Π -notation to emphasize the connection to functions.

anything of type **Type** and “terms” are things of type A where A has type **Type**.

We can use the typing rules to construct a typing derivation for the identity function as follows.

$$\frac{\frac{y : x \in (x : \mathbf{Type}, y : x)}{x : \mathbf{Type}, y : x \vdash y : x} \quad \frac{x : \mathbf{Type} \in (x : \mathbf{Type})}{x : \mathbf{Type} \vdash x : \mathbf{Type}}}{\frac{x : \mathbf{Type} \vdash \lambda y. y : (y : x) \rightarrow x \quad \vdash \mathbf{Type} : \mathbf{Type}}{\vdash \lambda x. \lambda y. y : (x : \mathbf{Type}) \rightarrow (y : x) \rightarrow x}}$$

This derivation uses rule T-LAMBDA to type check the two λ -expressions, before using the variable rule to ensure that both the body of the function x and its type y are well-typed. Finally, this rule also uses rule T-TYPE to show that **Type** itself is a valid type, see below.

More typing rules: Types Observe in the typing derivation above that the rule for typing lambda-expressions has second precondition: we need to make sure that when we add assumptions $x : A$ to the context, that A really is a type. Without this precondition, the rules would allow us to derive this nonsensical type for the polymorphic identity function.

$$\vdash \lambda x. \lambda y. y : (x : \mathbf{True}) \rightarrow (y : x) \rightarrow x$$

This precondition means that we need some rules that conclude that types are actually types. For example, the type of a function is itself a type, so we will declare it so with rule T-PI. This rule also ensures that the domain and range of the function are also types.

Likewise, for polymorphism we need the somewhat perplexing rule T-TYPE, that declares, by fiat, that **Type** is a valid **Type**.⁶

More typing rules: Applications The application rule, rule T-APP, requires that the type of the argument matches the domain type of the function. However, note that because the body of the function type B could have x free in it, we need also need to substitute the argument b for x in the result.

Example: applying the polymorphic identity function In `pi-forall` we should be able to apply the polymorphic identity function to itself. When we do this, we need to first provide the type of `id`, producing an expression of type $(y : ((x : \mathbf{Type}) \rightarrow (y : x) \rightarrow x)) \rightarrow ((x : \mathbf{Type}) \rightarrow (y : x) \rightarrow x)$. This function can take `id` as an argument.

```
idid : (x:Type) → (y : x) → x
idid = id ((x:Type) → (y : x) → x) id
```

⁶Note that this rule makes our language inconsistent as a logic, as it can encode Girard’s

Example: Church booleans Because we have (impredicative) polymorphism, we can *encode* familiar types, such as booleans. The idea behind this encoding, called a Church encoding, is to represent terms by their eliminators. In other words, what is important about the value `true`? The fact that when you get two choices, you pick the first one. Likewise, `false` “means” that with the same two choices, you should pick the second one. With parametric polymorphism, we can give the two terms the same type, which we’ll call `bool`.

```
bool : Type
bool = (x : Type) → x → x → x
```

```
true : bool
true = λx. λy. λz. y
```

```
false : bool
false = λx. λy. λz. z
```

Thus, a conditional expression just takes a boolean and returns it.

```
cond : bool → (x:Type) → x → x → x
cond = λb. b
```

Example: Logical “and” We can also encode a logical “and” data structure using a Church encoding.

```
and : Type → Type → Type
and = λp. λq. (c: Type) → (p → q → c) → c
```

```
conj : (p:Type) → (q:Type) → p → q → and p q
conj = λp.λq. λx.λy. λc. λf. f x y
```

```
proj1 : (p:Type) → (q:Type) → and p q → p
proj1 = λp. λq. λa. a p (λx.λy.x)
```

```
proj2 : (p:Type) → (q:Type) → and p q → q
proj2 = λp. λq. λa. a q (λx.λy.y)
```

```
and_commutates : (p:Type) → (q:Type) → and p q → and q p
and_commutates = λp. λq. λa. conj q p (proj2 p q a) (proj1 p q a)
```

4 From typing rules to a typing algorithm

The rules that we have developed so far are great for saying *what* terms should type check, but they don’t say *how* type checking works. We’ve developed these rules without thinking about how we would implement them.

A set of typing rules, or *type system*, is called *syntax-directed* if it is readily

paradox.

apparent how to interpret that collection of typing rules as code. In other words, for some type systems, we can directly translate them to some Haskell function.

For this type system, we would like to implement the following Haskell function, that when given a term and a typing context, represented as a list of pairs of variables and their types, produces the type of the term, if it exists.⁷

```
type Ctx = [(Var, Type)]
```

```
inferType :: Term → Ctx → Maybe Type
```

Let’s look at our rules in Figure 2. Is the definition of this function straightforward? For example, in the variable rule, as long as we can look up the type of a variable in the context, we can produce its type. That means that, assuming that there is some function `lookupTy`, with type `Ctx → Var → Maybe Type`, this rule corresponds to the following case of the `inferType` function. So far so good!

```
inferType (Var x) ctx = lookupTy ctx x
```

Likewise, the case of the typing function for the **Type** term is also straightforward. When we see the term **Type**, we know immediately that it is its own type.

```
inferType Type ctx = Just Type
```

The only stumbling block for the algorithm is rule T-LAMBDA. To type check a function, we need to type check its body when the context has been extended with the type of the argument. But, like Haskell, the type of the argument *A* is not annotated on the function in `pi-forall`. So where does it come from?

There is actually an easy fix to turn our current system into an algorithmic one. We could just annotate λ -expressions with the types of the abstracted variables. But, to be ready for future extension, we’ll do something else.

Look at our example `pi-forall` code above: the only types that we wrote were the types of definitions. It is good style to do that. Furthermore, there is enough information there for type checking—wherever we define a function, we can look at those types to know what type its argument should have. So, by changing our point of view, we can get away without annotating lambdas with those argument types.

4.1 A bidirectional type system

Let’s redefine the system using two judgments. The first one is similar to the judgement that we saw above, and we will call it *type inference*.⁸ This judgement will be paired with (and will depend on) a second judgement, called *type*

⁷Note: If you are looking at the `pi-forall` implementation, note that this is not the final type of `inferType`.

⁸The term *type inference* usually refers to much more sophisticated deduction of an expressions type, in the context of much less information encoded in the syntax of the language. We’re not doing anything difficult here, just noting that we can read the judgment with *A* as

$\boxed{\Gamma \vdash a \Rightarrow A}$		<i>(in context Γ, infer that term a has type A)</i>	
I-VAR $\frac{x : A \in \Gamma}{\Gamma \vdash x \Rightarrow A}$	I-APP-SIMPLE $\frac{\Gamma \vdash a \Rightarrow (x:A) \rightarrow B \quad \Gamma \vdash b \Leftarrow A}{\Gamma \vdash a \ b \Rightarrow B[b/x]}$	I-PI $\frac{\Gamma \vdash A \Leftarrow \mathbf{Type} \quad \Gamma, x:A \vdash B \Leftarrow \mathbf{Type}}{\Gamma \vdash (x:A) \rightarrow B \Rightarrow \mathbf{Type}}$	
I-TYPE $\frac{}{\Gamma \vdash \mathbf{Type} \Rightarrow \mathbf{Type}}$		I-ANNOT $\frac{\Gamma \vdash a \Leftarrow A}{\Gamma \vdash (a : A) \Rightarrow A}$	
$\boxed{\Gamma \vdash a \Leftarrow A}$		<i>(in context Γ, check that term a has type A)</i>	
C-LAMBDA $\frac{\Gamma, x:A \vdash a \Leftarrow B \quad \Gamma \vdash A \Leftarrow \mathbf{Type}}{\Gamma \vdash \lambda x.a \Leftarrow (x:A) \rightarrow B}$		C-INFER-SIMPLE $\frac{\Gamma \vdash a \Rightarrow A}{\Gamma \vdash a \Leftarrow A}$	

Figure 3: (Simple) Bidirectional type system

checking, that takes advantage of known type information, such as the annotations on top-level definitions.

We express these judgements using the notation defined in Figure 3 and implement them in Haskell using the mutually-recursive functions `inferType` and `checkType`. Furthermore, to keep track of which rule is in which judgement, rules that have inference as a conclusion start with I- and rules that have checking as a conclusion start with C-.

Let's compare these rules with our original typing rules. For rule I-VAR, we need to only change the colon to an inference arrow. The context tells us the type to infer.

On the other hand, in rule C-LAMBDA we should check λ -expressions against a known type. If that type is provided, we can propagate it to the body of the lambda expression. We also want to check that A is a **Type**.

The rule for applications, rule I-APP-SIMPLE, is in inference mode. Here, we first infer the type of the function, but once we have that type, we may use it to check the type of the argument. This mode change means that λ -expressions that are arguments to other functions (like `map`) do not need any annotation.

For types, it is apparent their type is **Type**, so rules I-PI and I-TYPE just continue to infer that.

Notice that this system is incomplete. There are inference rules for every form of expression except for lambda. On the other hand, only lambda expressions can be checked against types. We can make the checking judgement more

an output.

applicable with rule C-INFER-SIMPLE that allows us to use inference whenever a checking rule doesn't apply.

Now, let's think about the reverse problem a bit. There are programs that the checking system won't admit but would have been acceptable by our first system. What do they look like?

Well, they involve applications of explicit lambda terms:

$$\frac{\vdash \lambda x.x : \mathbf{Bool} \rightarrow \mathbf{Bool} \quad \vdash \mathbf{True} : \mathbf{Bool}}{\vdash (\lambda x.x) \mathbf{True} : \mathbf{Bool}} \text{ T-APP}$$

This term doesn't type check in the bidirectional system because application requires the function to have an inferable type, but lambdas don't. However, there is not that much need to write such terms. We can always replace them with something equivalent by doing a β -reduction of the application (in this case, just replace the term with **True**).

In fact, the bidirectional type system has the property that it only checks terms in *normal* form, i.e. those that do not contain any β -reductions.

Type annotations To type check nonnormal forms in **pi-forall**, we also add typing annotations as a new form of expression to **pi-forall**, written $(a : A)$, and add rule I-ANNOT to the type system.

Type annotations allow us to supply known type information anywhere, not just at top level. For example, we can construct this derivation.

$$\frac{\vdash (\lambda x.x : \mathbf{Bool} \rightarrow \mathbf{Bool}) \Rightarrow \mathbf{Bool} \rightarrow \mathbf{Bool} \quad \vdash \mathbf{True} \Leftarrow \mathbf{Bool}}{\vdash (\lambda x.x : \mathbf{Bool} \rightarrow \mathbf{Bool}) \mathbf{True} \Rightarrow \mathbf{Bool}} \text{ I-APP}$$

The nice thing about the bidirectional system is that it reduces the number of annotations that are necessary in programs that we want to write. As we will see, checking mode will be even more important as we add more terms to the language.

Furthermore, we want to convince ourselves that the bidirectional system checks the same terms as the original type system. This means that we want to prove a property like this one:

Lemma 4.1 (Correctness of the bidirectional type system). If $\Gamma \vdash a \Rightarrow A$ then $\Gamma \vdash a : A$. If $\Gamma \vdash a \Leftarrow A$ then $\Gamma \vdash a : A$.

On the other hand, the reverse property is not true. Even if there exists some typing derivation $\Gamma \vdash a : A$ for some term a , we may not be able to infer or check that term using the algorithm. However, all is not lost: there will always be some term a' that differs from a only in the addition of typing annotations that can be inferred or checked instead.

One issue with this bidirectional system is that it is not closed under substitution. What this means is that given some term b with a free variable, $\Gamma, x : A \vdash b \Leftarrow B$, and another term a with the same type $\Gamma \vdash a \Leftarrow A$ of that

variable, we *do not* have a derivation $\Gamma \vdash b[a/x] \Leftarrow A$. The reason is that types of variables are always inferred, but the term a , may need the type A to be supplied to the type checker. This issue is particularly annoying in rule I-APP when we replace a variable (inference mode) with a term that was validated in checking mode.

As a result, our type system infers types, but the types that are inferred do not have enough annotations to be checked themselves. We can express the property that does hold about our system, using this lemma:

Lemma 4.2 (Regularity). If $\Gamma \vdash a \Rightarrow A$ then $\Gamma \vdash A : \mathbf{Type}$.

This issue is not significant, but we could resolve it by adding an annotation before substitution. However, in our implementation of `pi-forall`, we do not do so to reduce clutter.

5 Putting it all together in a Haskell implementation

In the previous section, we defined a bidirectional type system for a small core language. Here, we'll start talking about what the implementation of this language looks like in Haskell.

First, an overview of the main files of the implementation. There are a few more source files in the repository (in the subdirectory `src/`), but these are the primary ones.

```
Syntax.hs      - specification of the AST
Parser.hs      - turn strings into AST
PrettyPrint.hs - displays AST in a (somewhat) readable form
Main.hs        - runs the parser and type checker

Environment.hs - defines the type checking monad
TypeCheck.hs   - implementation of the bidirectional type checker
```

5.1 Variable binding using the Unbound library [Syntax.hs]

One difficulty with implementing any sort of lambda calculus is the treatment of variable binding. Functions (λ -expressions) and their types (II-types) *bind* variables. In the implementation of our type checker, we'll need to be able to determine whether two terms are *alpha-equivalent*, calculate the *free variables* of a term, and perform *capture-avoiding substitution*. When we work with a λ -expression, we will want to be sure that the binding variable is *fresh*, that is, distinct from all other variables in the program or in the typing context.

In the `pi-forall` implementation, we use the `Unbound` library to get all of these operations (mostly) for free. This works because we use type constructors from this library in the definition of the abstract syntax of `pi-forall`, and those types specify the binding structure of `pi-forall` expressions.

First, the `Unbound` library defines a type for variable names, called `TName` and we use this type to define `TName`, the type of term names in our AST.

```
-- | Type used for variable names in Terms
type TName = Unbound.Name Term
```

The `Unbound.Name` type is indexed by `Term`, the AST type that it is a name for, whose definition we will see shortly. That way `Unbound` can make sure that we only substitute the right form of syntax for the right name, i.e. we can only replace a `TName` with a `Term`. The library includes an overloaded function `subst x a b`, which corresponds to the substitution we have been writing as $b[a/x]$ in our inference rules⁹, i.e. substitute a for x in b .

In general, we want to apply a substitution to many different sorts of syntax. For example, we may want to substitute a term a for a term name x in all of the terms that appear in the typing context. Or there may be other sorts of syntax that terms can be part of, and we want to be able to substitute for term variables through that syntax too. Therefore, `Unbound`'s substitution function has a type with the following general pattern for `subst a x b`.

```
class Subst a b where
  subst :: Name a → a → b → b
```

The `subst` function in this class ensures that when we see that `a` is the right sort of thing to stick in for `x`. The library can automatically generate instances of the `Subst` class.

With term names, we can define the abstract syntax that corresponds to our language above, using the following datatype.

```
data Term

= -- | type of types 'Type'
  Type

| -- | variables 'x'
  Var TName

| -- | abstractions 'λx. a'
  Lam (Unbound.Bind TName Term)

| -- | applications 'a b'
  App Term Term

| -- | function types '(x : A) → B'
  Pi Type (Unbound.Bind TName Type)

| -- | Annotated terms '( a : A )'
```

⁹See Guy Steele's talk about the notation for substitution [Ste17]. This is the most common

Ann Term Type

```
...  
deriving (Show, Generic)
```

As you can see, variables are represented by names. `Unbound`'s `Bind` type constructor declares the scope of the bound variables. Both `Lam` and `Pi` bind a single variable in a `Term`. However, in a `Pi` term, we also want to store the type `A`, the type of the bound variable `x`.

Because the syntax is all shared, a `Type` is just another name for a `Term`. We'll use this name in the source code just for documentation.

```
type Type = Term
```

Among other things, the `Alpha` class enables functions for alpha equivalence and free variable calculation, with the types shown below. Usually `Unbound` can derive instance of this class for us, so that we don't have to worry about defining these functions manually.

```
aeq :: Alpha a => a -> a -> Bool
```

For `Term`, we do not use the default definitions of the `Alpha` and `Subst` classes. Instead, we need to provide a bit more information. First, we would like the definition of alpha-equivalence of terms to ignore type annotations and source code positions for error messages. Therefore our instance does so for these constructors of the `Term` data type and defers to the generic definition for the rest.

For example, if `xName` and `yName` are two different names, then `Unbound`'s implementation of α -equivalence can equate terms that differ only in using these names for binding.

```
-- 'λx → x'  
idx :: Term  
idx = Lam (Unbound.bind xName (Var xName))  
  
-- 'λy → y'  
idy :: Term  
idy = Lam (Unbound.bind yName (Var yName))  
  
-- >>> Unbound.aeq idx idy  
-- True
```

Second, the `Subst` instance for `Term` requires telling `Unbound` where variables occur in the `Term`. This short function pattern matches to find the `Var` constructor, extracts its name and returns it to `Unbound`. We don't need to provide any more code to `Unbound` in the implementation of substitution; the generic implementation suffices.

```
instance Unbound.Subst Term Term where
```

mathematical notation for this operation.


```
isvar (Var x) = Just (SubstName x)
isvar _ = Nothing
```

For more information about `Unbound`, see the documentation for `unbound-generics`¹⁰.

5.2 A Type Checking monad [Environment.hs]

Recall that our plan is to write two mutually recursive functions for type checking of the following types:

```
inferType :: Term -> Ctx -> Maybe Type

checkType :: Term -> Type -> Ctx -> Bool
```

The inference function should take a term and a context and if it type checks, produce its type. The checking function should take a term, a context, and a type, and determine whether the term has that type. However, note that the types above are not very informative in the case of failure: if a term does not type check, then the result is `Nothing` or `False`.

For that reason, we'll do something a bit different. We'll define a *type checking monad*, called `TcMonad`, that will handle the plumbing for the typing context, and more importantly, allow us to return more information when a program doesn't type check. Therefore our two functions will have these types that hide the context and rely on throwing errors to indicate type checking failure.

```
inferType :: Term -> TcMonad Type

checkType :: Term -> Type -> TcMonad ()
```

Those of you who have worked with Haskell before may be familiar with the classes `MonadReader`, and the `MonadError`, which our type checking monad will be instances of. These two classes mean that that the `TcMonad` type can provide access to the type checking context and can throw errors. We'll wrap this functionality up a bit with the following functions. The first two encapsulate the sort of context manipulation we need to do: look up the types of variables and add a new assumption about the type of a variable to the context (both `Sig` and `Decl` are defined in `Syntax.hs`).

```
-- | Find the type of a name specified in the context
-- throwing an error if the name doesn't exist
lookupTy :: TName -> TcMonad Sig

-- | Extend the context with a new binding
extendCtx :: Decl -> TcMonad Term -> TcMonad Term

-- | Throw an error at the corrent location, assembling the error message
```

¹⁰<https://github.com/lambdageek/unbound-generics>

```

-- from a list of displayable objects. Terminates type checking
err  :: (Disp a) => [a] -> TcMonad b

-- | Print a warning, but continue the current computation.
warn :: (Disp a) => a -> TcMonad ()

```

We'll also need this monad to be a freshness monad, to support working with binding structure (i.e. a source of fresh names), and throw in `MonadIO` so that we can print warning messages.

5.3 Implementing the Type Checking Algorithm [Type-check.hs]

Now that we have the type checking monad available, we can start our implementation. For flexibility `inferType` and `checkType` will *both* be implemented by the same function, called `tcTerm`.

```

inferType :: Term -> TcMonad Type
inferType t = tcTerm t Nothing

checkType :: Term -> Type -> TcMonad ()
checkType tm ty = void $ tcTerm tm (Just ty)

```

The `tcTerm` function checks a term, producing its type. The second argument is `Nothing` in inference mode and an expected type in checking mode.

```
tcTerm :: Term -> Maybe Type -> TcMonad Type
```

We combine these two functions into a single definition so that we can share code between checking and inference mode. While none of the expression forms of the core language can take advantage of this functionality, it is useful for some extensions to the language.

The general structure of the `tcTerm` function starts with a pattern match for the various syntactic forms in inference mode:

```

tcTerm (Var x) Nothing = ...

tcTerm Type Nothing = ...

tcTerm (Pi tyA bnd) Nothing = ...

tcTerm (App t1 t2) Nothing = ...

```

Mixed in here, we also have a pattern for lambda expressions in checking mode:

```

tcTerm (Lam bnd) (Just (Pi tyA bnd2)) = ...    -- pass in the Pi type for the lambda expression
tcTerm (Lam _) (Just nf) = -- checking mode wrong type
  err [DS "Lambda expression has a function type, not ", DD nf]

```

There are also several cases for practical reasons (annotations, source code positions, etc.) and a few cases for homework.

Finally, the last case covers all other forms of checking mode, by calling inference mode and making sure that the inferred type is equal to the checked type. This case is the implementation of rule C-INFER.

```

tcTerm tm (Just ty) = do
  ty' <- inferType tm
  unless (Unbound.aeq ty' ty) $
    err [DS "Types don't match", DD ty, DS "and", DD ty']
  return ty

```

The function `aeq` ensures that the two types are alpha-equivalent. If they are not, it stops the computation and throws an error.

5.4 Exercise: Add Booleans and Σ -types

Some fairly standard typing rules for booleans are shown below.

$$\boxed{\Gamma \vdash a : A} \quad (\text{Booleans})$$

$$\begin{array}{c}
\text{T-BOOL} \\
\hline
\Gamma \vdash \mathbf{Bool} : \mathbf{Type}
\end{array}
\quad
\begin{array}{c}
\text{T-TRUE} \\
\hline
\Gamma \vdash \mathbf{True} : \mathbf{Bool}
\end{array}
\quad
\begin{array}{c}
\text{T-FALSE} \\
\hline
\Gamma \vdash \mathbf{False} : \mathbf{Bool}
\end{array}$$

$$\begin{array}{c}
\text{T-IF-SIMPLE} \\
\Gamma \vdash a : \mathbf{Bool} \quad \Gamma \vdash b_1 : A \quad \Gamma \vdash b_2 : A \\
\hline
\vdash \mathbf{if } a \mathbf{ then } b_1 \mathbf{ else } b_2 : A
\end{array}$$

Likewise, we can also extend the language with Σ -types, or products, where the type of the second component of the product can depend on the value of the first component.

$$\boxed{\Gamma \vdash a : A} \quad (\Sigma\text{-types})$$

$$\begin{array}{c}
\text{T-SIGMA} \\
\Gamma \vdash A : \mathbf{Type} \\
\Gamma, x : A \vdash B : \mathbf{Type} \\
\hline
\Gamma \vdash \{x : A \mid B\} : \mathbf{Type}
\end{array}
\quad
\begin{array}{c}
\text{T-PAIR} \\
\Gamma \vdash a : A \\
\Gamma \vdash b : B[a/x] \\
\hline
\Gamma \vdash (a, b) : \{x : A \mid B\}
\end{array}$$

$$\frac{\text{T-LETPAIR-WEAK} \quad \Gamma \vdash a : \{x:A_1 \mid A_2\} \quad \Gamma, x:A_1, y:A_2 \vdash b : B \quad \Gamma \vdash B : \mathbf{Type}}{\Gamma \vdash \mathbf{let} (x, y) = a \mathbf{in} b : B}$$

We deconstruct Σ -types using pattern matching. The simple rule for pattern matching introduces variables into the context when type checking the body of the `let` expression. These variables are not allowed to appear free in the result type of the pattern match.

This part of the homework has two parts:

1. First, on paper: rewrite the rules above in bidirectional style. Which rules should be inference rules? Which ones should be checking rules? If you are familiar with other systems, how do these rules compare?
2. Second, in Haskell: The code in `version1/` includes abstract and concrete syntax for booleans and Σ -types. The `pi-forall` file `version1/test/Hw1.pi` contains examples of using these new forms. Your job is to get this file to compile by filling in the missing cases in `version1/src/TypeCheck.hs` base on the bidirectional rules that you worked out in the first part of the exercise.

6 Equality in Dependently-Typed Languages

You may have noticed in the previous sections that there was something missing. Most of the examples that we did could have also been written in System F (or something similar)!

Next, we are going to think about how adding a notion of definitional equality can make our language more expressive.

The addition of definitional equality follows several steps. First, we will see why we even want this feature in the language in the first place. Next, we will create a declarative specification of definitional equality and extend our typing relation with a new rule that enables it to be used. After that, we'll talk about algorithmic versions of both the equality relation and how to introduce it into the algorithmic type system. Finally, we'll cover modifications to the Haskell implementation.

6.1 Motivating Example: Type level reduction

In full dependently-typed languages (and in full `pi-forall`) we can see the need for definitional equality. We want to equate types that are not merely alpha-equivalent, so that more expressions type check.

We saw yesterday an example where we wanted a definition of equality that was more expressive than alpha-equivalence. Recall our encoding for the logical and proposition:

```
and : Type → Type → Type
and = λp. λq. (c : Type) → (p → q → c) → c
```

Unfortunately, our definition of `conj` still doesn't type check:

```
conj : (p:Type) → (q:Type) → p → q → and p q
conj = λp.λq. λx.λy. λc. λf. f x y
```

Running this example with `version1` of the type checker produces the following error:

```
Checking module "Lec1"
Type Error:
../test/Lec1.pi:34:22:
  Function a should have a function type. Instead has type and p q
  When checking the term
    \p . \q . \a . a p ((\x . \y . x))
  against the signature
    (p : Type) -> (q : Type) -> (and p q) -> p
  In the expression
    a p ((\x . \y . x))
```

The problem is that even though we want `and p q` to be equal to the type `(c : Type) -> (p -> q -> c) -> c` the type checker does not treat these types as equal.

Note that the type checker already records in the environment that `and` is defined as `\p.\q. (c : Type) -> (p -> q -> c) -> c`. We'd like the type checker to look up this definition when it sees the variable `and` and beta-reduce this application.

6.2 Another example needing more expressive equality

As another example, in the full language, we might have a type of length indexed vectors, where vectors containing values of type `A` with length `n` can be given the type `Vec n A`. In this language we may have a safe head operation, that allows us to access the first element of the vector, as long as it is nonzero.

```
head : (A : Nat) → (n : Nat) → Vec (Succ n) A → Vec n A
head = ...
```

However, to call this function, we need to be able to show that the length of the argument vector is equal to `Succ n` for some `n`. This is ok if we know the length of the vector outright.

```
v1 : Vec 1 Bool
v1 = Cons [0] True Nil
```

So the application `head Bool 0 v1` will type check. (Note that `pi-forall` cannot infer the types `A` and `n`.)

However, if we construct the vector, its length may not be a literal natural number:

```
append : (n : Nat) → (m : Nat) → Vec m A → Vec n A → Vec (plus m n) A
append = ...
```

$$\boxed{\Gamma \vdash A = B} \quad (\text{Definitional Equality})$$

$$\begin{array}{c}
\text{E-BETA} \quad \frac{}{\Gamma \vdash (\lambda x. a) b = a[b/x]} \quad \text{E-REFL} \quad \frac{}{\Gamma \vdash A = A} \quad \text{E-SYM} \quad \frac{\Gamma \vdash A = B}{\Gamma \vdash B = A} \quad \text{E-TRANS} \quad \frac{\Gamma \vdash A_1 = A_2 \quad \Gamma \vdash A_2 = A_3}{\Gamma \vdash A_1 = A_3} \\
\\
\text{E-PI} \quad \frac{\Gamma \vdash A_1 = A_2 \quad \Gamma, x : A_1 \vdash B_1 = B_2}{\Gamma \vdash (x : A_1) \rightarrow B_1 = (x : A_2) \rightarrow B_2} \quad \text{E-LAM} \quad \frac{\Gamma, x : A_1 \vdash b_1 = b_2}{\Gamma \vdash \lambda x. b_1 = \lambda x. b_2} \\
\\
\text{E-APP} \quad \frac{\Gamma \vdash a_1 = a_2 \quad \Gamma \vdash b_1 = b_2}{\Gamma \vdash a_1 b_1 = a_2 b_2} \quad \text{E-LIFT} \quad \frac{\Gamma, x : A \vdash b : B \quad \Gamma \vdash a_1 = a_2}{\Gamma \vdash b[a_1/x] = b[a_2/x]} \quad \text{E-ANNOT} \quad \frac{\Gamma \vdash a_1 = a_2}{\Gamma \vdash (a_1 : A) = a_2}
\end{array}$$

Figure 4: Definitional equality for core `pi-forall`

In that case, to get `head Bool 1 (append v1 v1)` to type check, we need to show that the type `Vec (Succ 1) Bool` is equal to the type `Vec (plus 1 1) Bool`. If our definition of type equality is *alpha-equivalence*, then this equality will not hold. We need to enrich our definition of equality so that it equates more terms.

6.3 Defining equality

The main idea is that we will:

- establish a new judgement that defines when types are equal

$$\boxed{\Gamma \vdash A = B}$$

- add the following rule to our type system so that it works “up-to” our defined notion of type equivalence

$$\text{T-CONV} \quad \frac{\Gamma \vdash a : A \quad \Gamma \vdash A = B}{\Gamma \vdash a : B}$$

- Figure out how to revise the *algorithmic* version of our type system so that it supports the above rule.

What is a good definition of equality? We have implicitly started with a very simple one: identify terms up to α -equivalence. But we can do better.

The rules in Figure 4 define what it means for two terms to be equivalent, by stating the properties that we want to be true of this relation. Rule E-BETA ensures that our relation *contains beta-equivalence*. Terms that evaluate to each other should be equal. Rules E-REFL, E-SYM, and E-TRANS makes sure that this relation is an *equivalence relation*. Furthermore, we want this relation to be a *congruence relation* (i.e. if subterms are equal, then larger terms are equal), as specified by rules rules E-PI, E-LAM, and E-APP. We also want to be sure that this relation has “functionality” (i.e. we can lift equalities). We declare so with rule E-LIFT. Finally, we want to ignore type annotations, so rule E-ANNOT allows the equality judgment to skip over them on the left (and with the help of rule E-SYM, on the right as well).

(Note: if we add booleans and Σ -types, we will also need corresponding β -equivalence rules and congruence rules for those constructs.)

6.4 Using definitional equality in the algorithm

We would like to consider our type system as having the following rule:

$$\frac{\text{T-CONV} \quad \Gamma \vdash a : A \quad \Gamma \vdash A = B}{\Gamma \vdash a : B}$$

but this rule is not syntax-directed; it can be used at any place in a derivation so how do we know where to check types for equality? And which types should we check?

It turns out that in our bidirectional system, there are only a few places where type equality matters.

- We need to check for equality when we switch from checking mode to inference mode in the algorithm in rule rule C-INFER. Here we need to ensure that the type that we infer is the same as the type that is passed to the checker.

$$\frac{\text{C-INFER} \quad \Gamma \vdash a \Rightarrow A \quad \Gamma \vdash A \Leftrightarrow B}{\Gamma \vdash a \Leftarrow B}$$

In this case, the types A and B are available to the type checker, and the our equality algorithm must decide whether they are equal. We use $\Gamma \vdash A \Leftrightarrow B$ as the notation for an algorithmic version of our equality relation—the relation we defined in Figure 4 does not readily lead to an algorithm. (More on this algorithmic version in Section 6.5 below.)

- In the rule for application, when we infer the type of the function we need to make sure that the function actually has a function type. But we don’t really know what the domain and co-domain of the function should be. We’d like our algorithm for type equality to be able to figure this out for us.

$$\boxed{\mathbf{whnf} \ a \rightsquigarrow v} \quad (\text{Weak head normal form reduction})$$

$$\begin{array}{c}
\text{WHNF-LAM-BETA} \\
\frac{\mathbf{whnf} \ a \rightsquigarrow (\lambda x. a') \quad \mathbf{whnf} \ (a'[b/x]) \rightsquigarrow v}{\mathbf{whnf} \ a \ b \rightsquigarrow v}
\end{array}
\quad
\begin{array}{c}
\text{WHNF-LET-BETA} \\
\frac{\mathbf{whnf} \ (b[a/x]) \rightsquigarrow v}{\mathbf{whnf} \ (\mathbf{let} \ x = a \ \mathbf{in} \ b) \rightsquigarrow v}
\end{array}$$

$$\begin{array}{c}
\text{WHNF-TYPE} \\
\frac{}{\mathbf{whnf} \ \mathbf{Type} \rightsquigarrow \mathbf{Type}}
\end{array}
\quad
\begin{array}{c}
\text{WHNF-LAM} \\
\frac{}{\mathbf{whnf} \ \lambda x. a \rightsquigarrow \lambda x. a}
\end{array}$$

$$\begin{array}{c}
\text{WHNF-PI} \\
\frac{}{\mathbf{whnf} \ (x:A) \rightarrow B \rightsquigarrow (x:A) \rightarrow B}
\end{array}
\quad
\begin{array}{c}
\text{WHNF-VAR} \\
\frac{}{\mathbf{whnf} \ x \rightsquigarrow x}
\end{array}
\quad
\begin{array}{c}
\text{WHNF-LAM-CONG} \\
\frac{\mathbf{whnf} \ a \rightsquigarrow \mathit{path}}{\mathbf{whnf} \ a \ b \rightsquigarrow \mathit{path} \ b}
\end{array}$$

Figure 5: Relation computing the weak head normal form of a term

$$\frac{\text{I-APP} \quad \Gamma \vdash a \Rightarrow A \quad \mathbf{whnf} \ A \rightsquigarrow (x:A_1) \rightarrow B \quad \Gamma \vdash b \Leftarrow A_1}{\Gamma \vdash a \ b \Rightarrow B[b/x]}$$

In this case, we are given a single type A and we need to know whether it is equivalent to some Π -type. Because Π -types are *head* forms, we can do this via reduction. This rule evaluates the type A to its head form, using the rules shown in Figure 5. If that form is a Π -type, then we can access its subcomponents.

- Whenever we call `checkType` we should call it with a term that has already been reduced to normal form. This will allow rule C-LAM to match against a literal function type.

Weak head normal form The following judgement (shown in Figure 5)

$$\boxed{\mathbf{whnf} \ a \rightsquigarrow v}$$

describes when a term a reduces to some result v in *weak head normal form* (*whnf*). For closed terms, these rules correspond to a big-step evaluation relation and produce a value. For open terms, the rules produce a term that is either headed by a variable or by some head form (i.e. λ or Π).

For example, the term

$$(\lambda \mathbf{x}. \mathbf{x}) \ (\lambda \mathbf{x}. \mathbf{x})$$

is not in *whnf*, because there is more reduction to go to get to the head. On the other hand, even though there are still internal reductions possible, the terms

`λy. (λx.x) (λx.x)`

and

`(y:Type) → (λx.x) Bool`

are in weak head normal form. Furthermore, the term `x y` is also in weak head normal form, because, even though we don't know what the head form is, we cannot reduce the term any more.

Because evaluation may not terminate, this relation is semi-decidable. However, it is syntax-directed so we can readily express it as a Haskell function.

Using algorithmic equality In the `pi-forall` implementation, the function that corresponds to $\Gamma \vdash a \Leftrightarrow b$ has type

```
equate :: Term -> Term -> TcMonad ()
```

This function ensures that the two provided types are equal, or throws a type error if they are not.

Additionally, the `pi-forall` includes the function

```
ensurePi :: Type -> TcMonad (TName, Type, Type)
```

that checks the given type to see if the given is equal to some Π -type of the form $(x:A_1) \rightarrow A_2$, and if so returns `x`, `A1` and `A2`.

This function is defined in terms of a helper function that implements the rules shown in Figure 5.

```
whnf :: Term -> TcMonad Term
```

In `version2` of the implementation, these functions are called in a few places:

- `equate` is called at the end of `tcTerm`
- `ensurePi` is called in the `App` case of `tcTerm`
- `whnf` is called in `checkType`, before the call to `tcTerm`, to make sure that we are using the head form in checking mode.

6.5 Implementing definitional equality (see `Equal.hs`)

The rules for $\Gamma \vdash a = b$ *specify* when terms should be equal, but they are not an algorithm. But how do we implement its algorithmic analogue $\Gamma \vdash a \Leftrightarrow b$?

There are several ways to do so.

The easiest one to explain is based on reduction—for `equate` to reduce the two arguments to some normal form and then compare those normal forms for equivalence.

One way to do this is with the following algorithm:

```
equate t1 t2 = do
```

```

nf1 <- reduce t1 -- full reduction, not just weak head reduction
nf2 <- reduce t2
Unbound.aeq nf1 nf2

```

However, we can do better. Sometimes we can equate the terms without completely reducing them. Because reduction can be expensive (or even non-terminating) we'd like to only reduce as much as necessary.

Therefore, the implementation of `equate` has the following form.

```

equate t1 t2 = do
  if (Unbound.aeq t1 t1) then return () else do
    nf1 <- whnf t1 -- reduce only to 'weak head normal form'
    nf2 <- whnf t2
    case (nf1,nf2) of
      (Lam bnd1, Lam bnd2) -> do
        -- ignore variable names, but use the same
        -- fresh name for both lambda bodies
        (_, b1, _, b2) <- Unbound.unbind2Plus bnd1 bnd2
        equate b1 b2
      (App a1 a2, App b1 b2) -> do
        equate a1 b1
        equateArg a2 b2

    ... -- all other matching head forms

    -- | head forms don't match throw an error
    (_,_) -> Env.err ...

```

Above, we first check whether the terms are already alpha-equivalent. If they are, we do not need to do anything else, we can return immediately (i.e. without throwing an error, which is what the function does when it decides that the terms are not equal). The next step is to reduce both terms to their weak head normal forms. If two terms have different head forms, then we know that they must be different terms so we can throw an error. If they have the same head forms, then we can call the function recursively on analogous subcomponents until the function either terminates or finds a mismatch.

Why do we use weak head reduction vs. full reduction?

- We can implement deferred substitutions for variables. Note that when comparing terms we need to have their definitions available. That way we can compute that `(plus 3 1)` weak head normalizes to 4, by looking up the definition of `plus` when needed. However, we don't want to substitute all variables through eagerly—not only does this make extra work, but error messages can be extremely long.
- Furthermore, we allow recursive definitions in `pi-forall`, so normalization may just fail completely, if we unfold recursive definitions inside functions before they are applied. In contrast, the definition based on `whnf`

only unfolds recursive definitions when they stand in the way of equivalence, so avoids some infinite loops in the type checker.

Note that we don't have a complete treatment of equality. There will always be terms that can cause `equate` to loop forever.

7 Dependent pattern matching and propositional equality

One of the powerful features of dependently-typed languages is the idea of “learning by testing”. In other words, the information gained in pattern matching can be reflected into the type system, allowing it to reason in a flow-sensitive manner. We'll demonstrate how this works in this section.

7.1 More expressive rules for if and Σ -type elimination

Consider our elimination rules for if:

$$\frac{\text{T-IF-SIMPLE} \quad \Gamma \vdash a : \mathbf{Bool} \quad \Gamma \vdash b_1 : A \quad \Gamma \vdash b_2 : A}{\vdash \text{if } a \text{ then } b_1 \text{ else } b_2 : A}$$

We can strengthen this rule by making the result type A depend on whether the condition is true or false when type checking each branch. In this way, our type system is flow-sensitive: we do different things while type checking based on how control flows in a `pi-forall` program.

$$\frac{\text{T-IF-FULL} \quad \Gamma \vdash a : \mathbf{Bool} \quad \Gamma \vdash b_1 : A[\mathbf{True}/x] \quad \Gamma \vdash b_2 : A[\mathbf{False}/x]}{\Gamma \vdash \text{if } a \text{ then } b_1 \text{ else } b_2 : A[a/x]}$$

For example, here is a simple program that requires this stronger rule to type check. The function `T` is an example of a *large elimination*: a function that destructs a boolean value to produce a type. Then in the type of `bar` below, each branch can return a different type of result, according to `T`.

```
-- function from booleans to types
T : Bool → Type
T = λb. if b then Unit else Bool

-- returns Unit when the argument is True
bar : (b : Bool) → T b
bar = λb . if b then () else True
```

It turns out that this strong typechecking rule is difficult to implement. In fact, it is not syntax-directed because A and x are not fixed by the syntax. Given $A[\mathbf{True}/x]$ and $A[\mathbf{False}/x]$ and $A[a/x]$ (or anything that they are definitionally equal to!) how can we figure out whether they correspond to each other? Neither inference nor checking mode works, because we don't ever see the type A with

just the variable x .

So, we'll not be so ambitious in `pi-forall`. We'll only allow this refinement when the scrutinee is a variable, deferring to the weaker, non-refining typing rule for all other cases.

$$\frac{\text{T-IF} \quad \Gamma \vdash x : \mathbf{Bool} \quad \Gamma \vdash b_1 : A[\mathbf{True}/x] \quad \Gamma \vdash b_2 : A[\mathbf{False}/x]}{\Gamma \vdash \mathbf{if } x \mathbf{ then } b_1 \mathbf{ else } b_2 : A}$$

And, in the our bidirectional system, we'll only allow refinement when we are in checking mode. That way, because a is just x , we can easily identify how the type should vary in each branch.

$$\frac{\text{C-IF} \quad \Gamma \vdash x \Rightarrow \mathbf{Bool} \quad \Gamma \vdash b_1 \Leftarrow A[\mathbf{True}/x] \quad \Gamma \vdash b_2 \Leftarrow A[\mathbf{False}/x]}{\Gamma \vdash \mathbf{if } x \mathbf{ then } b_1 \mathbf{ else } b_2 \Leftarrow A}$$

To implement this rule, we need only remember that x is `True` or `False` when checking the individual branches of the if expression.¹¹

We can modify the elimination rule for Σ -types similarly. Here, we are not learning which branch was taken, but we are still learning something. If we get to the body of the `let` expression, we know that the scrutinee terminated and produced some product value.

$$\frac{\text{C-LETPAIR} \quad \Gamma \vdash z \Rightarrow \{x : A_1 \mid A_2\} \quad \Gamma, x : A_1, y : B_2 \vdash b \Leftarrow B[(x, y)/z] \quad \Gamma \vdash B \Leftarrow \mathbf{Type}}{\Gamma \vdash \mathbf{let } (x, y) = z \mathbf{ in } b \Leftarrow B}$$

This modification changes our definition of Σ -types from weak Σ s to strong Σ -types. With either typing rule, we can define the first projection.

```
fst : (A : Type) → (B : A → Type) → (p : { x2 : A | B x2 }) → A
fst = λA B p. let (x,y) = p in x
```

But, weak Σ -types cannot define the second projection. The following code only type checks when the above rule is available.

¹¹Here is an alternative version, for inference mode only, suggested during a prior summer school:

$$\frac{\text{I-IF-ALT} \quad \Gamma \vdash a \Rightarrow \mathbf{Bool} \quad \Gamma \vdash b_1 \Rightarrow B_1 \quad \Gamma \vdash b_2 \Rightarrow B_2}{\Gamma \vdash \mathbf{if } a \mathbf{ then } b_1 \mathbf{ else } b_2 \Rightarrow \mathbf{if } a \mathbf{ then } B_1 \mathbf{ else } B_2}$$

It has a nice symmetry—if expressions are typed by if types. Note however, to make this rule work, we'll need a stronger definitional equivalence than we have. In particular, we'll want our definition of equivalence to support the following equality:

$$\frac{\text{E-IF-ETA}}{\Gamma \vdash \mathbf{if } a \mathbf{ then } b \mathbf{ else } b = b}$$

That way, if the type of the two branches of the if does not actually depend on the boolean

$$\boxed{\Gamma \vdash a : A} \quad (Typing)$$

$$\begin{array}{c}
\text{T-REFL} \\
\frac{\Gamma \vdash a = b}{\Gamma \vdash \mathbf{refl} : a = b}
\end{array}
\quad
\begin{array}{c}
\text{T-EQ} \\
\frac{\Gamma \vdash a : A \quad \Gamma \vdash b : A}{\Gamma \vdash a = b : \mathbf{Type}}
\end{array}$$

$$\begin{array}{c}
\text{T-SUBST} \\
\frac{\Gamma \vdash a : A[a_1/x][\mathbf{refl}/y] \quad \Gamma \vdash b : a_1 = a_2}{\Gamma \vdash \mathbf{subst } a \text{ by } b : A[a_2/x][b/y]}
\end{array}
\quad
\begin{array}{c}
\text{T-CONTRA} \\
\frac{\Gamma \vdash a : \mathbf{True} = \mathbf{False}}{\Gamma \vdash \mathbf{contra } a : A}
\end{array}$$

Figure 6: Propositional equality (Specification rules)

```

snd : (A : Type) → (B : A → Type) → (p : { x2 : A | B x2 }) → B (fst A B p)
snd = λA B p. let (x,y) = p in y

```

(Try this out using `version2` of the implementation and the `Lec2.pi` input file.)

7.2 Propositional equality

You started proving things right away in Coq or Agda with an equality proposition. For example, in Coq, when you say

```
Theorem plus_0_n : forall n : nat, 0 + n = n.
```

you are using a built in type, `a = b` that represents the proposition that two terms are equal.

As a step towards more general indexed datatypes, we'll start by adding a propositional equality type to `pi-forall`.

The main idea of the equality type is that it converts a *judgement* that two types are equal into a *type* that evaluates to a value only when two types are equal.¹² In other words, the new value form `refl` has type `a = b`, as shown in rule T-REFL in Figure 6.

Sometimes, you might see rule T-REFL written as follows:

$$\begin{array}{c}
\text{T-REFL-ALT} \\
\hline
\Gamma \vdash \mathbf{refl} : a = a
\end{array}$$

However, this rule is equivalent to the above version because types are equal up to congruence. If we know that $\vdash a = b$, then we know that the type `a = a` is equal to the type `a = b`. In a type system with conversion, it doesn't matter which of these equal types we choose for the conclusion of the rule.

value, we can convert the if expression into a more useful type.

¹²Recall that all types are inhabited by an infinite loop in `pi-forall`.

$$\boxed{\Gamma \vdash a \Leftarrow A} \quad (Checking)$$

$$\begin{array}{c}
\text{C-SUBST-LEFT} \quad \Gamma \vdash y \Rightarrow B \\
\text{C-SUBST-RIGHT} \quad \Gamma \vdash y \Rightarrow B \\
\text{C-REFL} \quad \Gamma \vdash a \Leftrightarrow b \\
\hline
\Gamma \vdash \mathbf{refl} \Leftarrow a = b
\end{array}$$

$$\begin{array}{c}
\text{C-SUBST-LEFT} \quad \Gamma \vdash y \Rightarrow B \\
\text{C-SUBST-RIGHT} \quad \Gamma \vdash y \Rightarrow B \\
\text{C-REFL} \quad \Gamma \vdash a \Leftrightarrow b \\
\hline
\Gamma \vdash \mathbf{subst} \, a \, \mathbf{by} \, y \Leftarrow A
\end{array}$$

$$\begin{array}{c}
\text{C-CONTRA} \\
\Gamma \vdash a : A \quad \mathbf{whnf} \, A \rightsquigarrow (a_1 = a_2) \\
\mathbf{whnf} \, a_1 \rightsquigarrow \mathbf{True} \quad \mathbf{whnf} \, a_2 \rightsquigarrow \mathbf{False} \\
\hline
\Gamma \vdash \mathbf{contra} \, a \Leftarrow B
\end{array}$$

$$\boxed{\Gamma \vdash a \Rightarrow A} \quad (Inference)$$

$$\begin{array}{c}
\text{I-EQ} \\
\Gamma \vdash a \Rightarrow A \quad \Gamma \vdash b \Rightarrow B \quad \Gamma \vdash A \Leftrightarrow B \\
\hline
\Gamma \vdash (a = b) \Rightarrow \mathbf{Type}
\end{array}$$

Figure 7: Propositional equality (Bidirectional rules)

However, in the bidirectional version of the rules, shown in Figure 7, it does make a difference. The algorithmic typing rule for **refl** requires checking mode, so that we know which types we need to show equivalent. (See rule C-REFL.) Furthermore, this rule allows a and b to differ so that there are no restrictions on its use. Instead, to implement this rule, we need only compare a and b for definitional equality, as we did in Section 6.

An equality type is well-formed when the terms on both sides of the equality have the same type. In other words, when it implements *homogeneous* equality, as shown in rule T-EQ. In the algorithm, rule I-EQ, we infer the types of each side of the equality type and then check to make sure that those types are equal.

The elimination rule for propositional equality is written **subst** a **by** b in **pi-forall**. This term allows us to convert the type of one expression to another, as shown in rule T-SUBST-SIMPLE below. In this rule, we can change the type of some expression a by replacing an occurrence of some term a_1 with an equivalent type.

$$\begin{array}{c}
\text{T-SUBST-SIMPLE} \\
\Gamma \vdash a : A[a_1/x] \quad \Gamma \vdash b : a_1 = a_2 \\
\hline
\Gamma \vdash \mathbf{subst} \, a \, \mathbf{by} \, b : A[a_2/x]
\end{array}$$

How can we implement this rule? For simplicity, we'll play the same trick that we did with booleans, requiring that one of the sides of the equality be a variable.

$$\begin{array}{c}
\text{C-SUBST-LEFT-SIMPLE} \\
\frac{\Gamma \vdash b \Rightarrow B \quad \text{whnf } B \rightsquigarrow (x = a_2) \quad \Gamma \vdash a \Leftarrow A[a_2/x]}{\Gamma \vdash \text{subst } a \text{ by } b \Leftarrow A} \\
\\
\text{C-SUBST-RIGHT-SIMPLE} \\
\frac{\Gamma \vdash b \Rightarrow B \quad \text{whnf } B \rightsquigarrow (a_1 = x) \quad \Gamma \vdash a \Leftarrow A[a_1/x]}{\Gamma \vdash \text{subst } a \text{ by } b \Leftarrow A}
\end{array}$$

Note that this elimination form for equality is powerful. We can use it to show that propositional equality is symmetric and transitive.

```

sym : (A:Type) → (x:A) → (y:A) → (x = y) → y = x
trans : (A:Type) → (x:A) → (y:A) → (z:A)
        → (x = z) → (z = y) → (x = y)

```

Furthermore, we can also extend the elimination form for propositional equality with dependent pattern matching as we did for booleans. This corresponds to the elimination rule rule T-SUBST in Figure 6, that observes that the only way to construct a proof of equality is with the term **refl**.¹³

As above, the rule C-SUBST-RIGHT, shown in Figure 7, (and the corresponding left rule) only applies when b is also a variable.

One last addition: **contra**. If we can somehow prove a false statement, then we should be able to prove anything. A contradiction is a proposition between two terms that have different head forms. For now, we'll use rule T-CONTRA, shown in Figure 6, that requires a proof that **True** equals **False**. The algorithmic version of this rule, rule C-CONTRA, must use weak-head normalization multiple times: first to find the equality type and then to show that the two sides of it have different head forms. This rule must be in checking mode because we can use **contra** to inhabit any type.

7.2.1 Homework (pi-forall: equality)

Complete the file `Hw2.pi`. This file gives you practice with working with equality propositions in `pi-forall`.

8 Irrelevance: the \forall of pi-forall

Now, let's talk about irrelevance. In dependently typed languages, some arguments are “ghost” or “specificational” and only there for proofs. For efficient executables, we don't want to have to “run” these arguments, nor do we want them taking up space in data structures.

Functional languages do this all the time: they erase *type annotations* and *type arguments* as part of the compilation pipeline. This erasure makes sense

¹³This version of subst is similar to an eliminator for propositional equality called J. In our setting, this eliminator is very strong and can be used to derive “axiom” K or the uniqueness of identity proofs. The eliminator for propositional equality in Coq and Agda is not this expressive, leaving the type system open to extension with other proofs of propositional equality.

because of parametric polymorphic functions are not allowed to depend on types. The behavior of `map` must be the same no matter whether it is operating on a list of integers or a list of booleans.

In a dependently-typed language we'd like to erase types too. And erase proofs that are only there to make things type check. Coq does this by making a distinction between `Prop` and `Set`. Everything in `Set` stays around until run time, and is guaranteed not to depend on `Prop`.

We'll take another approach.

In `pi-forall` we have new kind of quantification, called "forall", that marks erasable arguments. We mark forall quantified arguments with square brackets. For example, we can mark the type argument of the polymorphic identity function as erasable.

```
id : [x:Type] → (y : x) → x
id = λ[x] y. y
```

When we apply such functions, we'll put the argument in brackets too, so we remember that `id` is not really using that type.

```
t = id [Bool] True
```

However, we need to make sure that irrelevant arguments really are irrelevant. We wouldn't want to allow this definition:

```
id' : [x:Type] → [y:x] → x
id' = λ[x][y]. y
```

Here `id'` claims that its second parameter `y` is erasable, but it is not. And, when we check this code with `version3` of `pi-forall`, it is rejected with the following error message.

```
Type Error:
Lec3.pi:16:16:
  Cannot access irrelevant variables in this context
  When checking the term
    \[x] [y]. y
  against the signature
    id' : [x : Type] -> [x] -> x
  In the expression
    y
```

However, note that casting is a relevant use of propositional equality. In the example below, `pi-forall` will prevent us from marking the argument `pf` as irrelevant.

```
proprel : [a : Type] → (pf : a = Bool) → (x : a) → Bool
proprel = λ[a] pf x .
  subst x by pf
```

The reason for this restriction is because the language includes non-termination. We don't know whether this proof is `Ref1` or some infinite loop. So this argument must be evaluated to be sure that the two types are actually equal. If

(somehow) we knew that the argument would always evaluate to `Ref1` we could erase it.

8.1 How do we rule this out in the type system?

We need to make sure that an irrelevant variable is not “used” in the relevant parts of the body of a λ -abstraction. How can we do so?

The key idea is that we mark variables in the context with an ϵ annotation, which is either $+$ (relevant) or $-$ (irrelevant). “Normal” variables, introduced by λ -expressions, will always be marked as relevant, and only relevant variables can be used in terms. As a result, we revise the variable typing rule to require the relevant annotation. In the `version3` implementation, type declarations in the environment now include an `Epsilon` tag, and the function `checkStage` compares this tag

$$\frac{\text{T-EVAR} \quad x :^+ A \in \Gamma}{\Gamma \vdash x : A}$$

An irrelevant abstraction introduces its variable into the context tagged with the irrelevant annotation. Because of the $-$ tag, such variables are inaccessible in the body of the function.

$$\frac{\text{T-ELAMBDA} \quad \Gamma, x :^- A \vdash a : B \quad \Gamma^+ \vdash A : \mathbf{Type}}{\Gamma \vdash \lambda[x].a : [x : A] \rightarrow B}$$

However, this variable should be available for use in *types* and other *irrelevant* parts of the term. To enable this use, we use a special context operation Γ^- , as in the second premise in rule T-ELAMBDA, and in the rule T-EAPP rule below.

$$\frac{\text{T-EAPP} \quad \Gamma \vdash a : [x : A] \rightarrow B \quad \Gamma^+ \vdash b : A}{\Gamma \vdash a[b] : B[b/x]}$$

This operation on the context, called *resurrection*, converts all $-$ tags on variables to be $+$. It represents a shift in our perspective: because we have entered an irrelevant part of the term, the variables that were not visible before are now available after this context modification.

Finally, when checking irrelevant Π types, we mark the variable with $+$ when we add it to the context so that it may be used in the range of the Π type. Otherwise, we would not be able to verify the type of the polymorphic identity function above.

$$\frac{\text{T-EPI} \quad \Gamma \vdash A : \mathbf{Type} \quad \Gamma, x :^+ A \vdash B : \mathbf{Type}}{\Gamma \vdash [x : A] \rightarrow B : \mathbf{Type}}$$

8.2 Compile-time irrelevance

What does checking irrelevance buy us? Because we know that irrelevant arguments are not actually used by the function, we know that they can be *erased* prior to execution.

However, there is another benefit to marking some arguments as irrelevant: we can ignore them during equivalence checking. When deciding whether two terms are equal, we don't need to look at their irrelevant components. In other words, we can add the following equality rule to our definitional equality in the case of applications with irrelevant arguments.

$$\frac{\text{E-EAPP} \quad \Gamma \vdash a_1 = a_2}{\Gamma \vdash a_1[b_1] = a_2[b_2]}$$

If you look at **version3** of the implementation, you can see where we use this idea in the definition of equality. Note in this version that arguments are now tagged with their relevance, and when we compare arguments that are tagged as irrelevant, we do nothing. In fact, what we are doing is actually *defining* equality over just the computationally relevant parts of the term instead of the entire term. And we've been doing this in a limited form from the beginning—recall that **version1** of `pi-forall` already ignores type annotations. Just as type annotations do not affect computation and can be ignored, the same reasoning holds for irrelevant arguments.

Why is compile-time irrelevance important?

- faster comparison: don't have to look at the whole term when comparing for equality
- more expressive: these parts of the term don't actually have to be equal

For example, consider the example below. The parameter p takes an irrelevant argument, so we know that it must be a constant function. Therefore it is sound to equate any two applications of p because we know that we will always get the same result.

```
irrelevance : (p : [i : Nat] → Type) → p [1] = p [2]
irrelevance = λp . Refl
```

9 Datatypes and Indexed Datatypes

Finally, we'll add datatypes with erasable arguments to `pi-forall`. The code to look at is the “complete” implementation in `full`.

Unfortunately, datatypes are both:

- Really important (you see them *everywhere* when working with languages like Coq, Agda, Idris, etc.)

- Really complicated (there are a *lot* of details). In general, datatypes subsume booleans, Σ -types, propositional equality, and can carry irrelevant arguments. So they combine all of the complexity of the previous sections in one construct.

Unlike the previous sections, where we could walk through all of the details of the specification of the type system, not to mention its implementation, we won't be able to do that here. There is just too much! The goal of this part is to give you enough information so that you can pick up the Haskell code and understand what is going on.

Even then, realize that the implementation that I'm giving you is not the complete story! Recall that we're not considering termination. That means that we can think about eliminating datatypes merely by writing recursive functions; without having to reason about whether those functions terminate. Coq, Agda and Idris include a lot of machinery for this termination analysis, and we won't cover any of it.

We'll work up the general specification of datatypes piece-by-piece, generalizing from features that we already know to more difficult cases. We'll start with "simple" datatypes, and then extend them with both parameters and indices.

9.1 “Dirt simple” datatypes

Our first goal is simple. What do we need to get the simplest examples of non-recursive and recursive datatypes working? By this I mean datatypes that you might see in Haskell or ML, such as `Bool`, `Void` and `Nat`.

9.1.1 Booleans

For example, one homework assignment was to implement booleans. Once we have booleans then we can

```
data Bool : Type where
  True
  False
```

In the homework assignment, we used `if` as the elimination form for boolean values.

```
not : Bool → Bool
not = λb . if b then False else True
```

For uniformity, we'll have a common elimination form for all datatypes, called **case** that has branches for all cases. (We'll steal Haskell syntax for case expressions, including layout.) For example, we might rewrite `not` with case like this:

```
not : Bool → Bool
not = λb .
  case b of
    True  → False
    False → True
```

9.1.2 Void

The simplest datatype of all is one that has no constructors!

```
data Void : Type where {}
```

Because there are no constructors, the elimination form for values of this type doesn't need any cases!

```
false_elim : (A:Type) → Void → A
false_elim = λA v . case v of {}
```

Void brings up the issue of *exhaustiveness* in case analysis. Can we tell whether there are enough patterns so that all of the cases are covered? This is something that our implementation should be able to do.

9.1.3 Nat

Natural numbers include a data constructor with an argument. For simplicity in the parser, those parentheses must be there.

```
data Nat : Type where
  Zero
  Succ of (Nat)
```

In case analysis, we can give a name to that argument in the pattern.

```
is_zero : Nat → Bool
is_zero = λx . case x of
  Zero → True
  Succ n → False
```

9.1.4 Dependently-typed data constructor args

Now, even in our “dirt simple” system, we’ll be able to encode some new structures, beyond what is available in functional programming languages like Haskell and ML. These structures won’t be all that useful yet, but as we add parameters and indices to our datatypes, they will be. For example, here’s an example of a datatype declaration where the data constructors have dependent types.

```
data SillyBool : Type where
  ImTrue of (b : Bool) (_ : b = True)
  ImFalse of (b: Bool) (_ : b = False)
```

9.2 Implementing the type checker for “dirt simple” datatypes

Datatype declarations, such as `data Bool`, `data Void` or `data Nat` extend the context with new type constants (aka type constructors) and new data constructors. It is as if we had added a bunch of new typing rules to the type system, such as:

$\boxed{\Gamma \vdash a : A}$ (Typing)

$\frac{\text{T-NAT}}{\Gamma \vdash \mathbf{Nat} : \mathbf{Type}}$	$\frac{\text{T-VOID}}{\Gamma \vdash \mathbf{Void} : \mathbf{Type}}$	$\frac{\text{T-ZERO}}{\Gamma \vdash \mathbf{Zero} : \mathbf{Nat}}$
$\frac{\text{T-SUCC} \quad \Gamma \vdash n : \mathbf{Nat}}{\Gamma \vdash \mathbf{Succ} \, n : \mathbf{Nat}}$	$\frac{\text{T-IMTRUE} \quad \begin{array}{l} \Gamma \vdash a_1 : \mathbf{Bool} \\ \Gamma \vdash a_2 : a_1 = \mathbf{True} \end{array}}{\Gamma \vdash \mathbf{ImTrue} \, a_1 \, a_2 : \mathbf{SillyBool}}$	

In the general form, a *simple* data type declaration includes a name and a list of data constructors.

```
data T : Type where
  K1          -- no arguments
  K2 of (A)    -- single arg of type A
  K3 of (x:A)  -- also single arg of type A, called x for fun
  K4 of (x:A)(y:B) -- two args, the type of B can mention A
  K5 of (x:A)[x = a] -- one arg, plus constraint about that arg
```

In fact, each data constructor takes a special sort of list of arguments that we call a *telescope*. (The word telescope for this structure was coined by de Bruijn [dB72] to describe the scoping behavior of this structure. The scope of each variable overlaps all of the subsequent ones, nesting like an expandable telescope.)

We can represent this structure in our implementation by adding a new form of declaration (some parts have been elided compared to `full`, so we’re building up to that version.)

```
-- | type constructor names
type TCName = String

-- | data constructor names
type DCName = String

data Decl =
  -- | Declaration for the type of a term ‘x : A’
  TypeSig Sig
  | -- | The definition of a particular name ‘x = a’
  Def TName Term
  ...

  | -- | Datatype definition ‘data T = ...’
  Data TCName [ConstructorDef]

-- | A Data constructor has a name and a telescope of arguments
data ConstructorDef = ConstructorDef DCName Telescope
```

$\Gamma \vdash \bar{a} \Leftarrow \Delta$	tcArgTele	Type check a list of arguments against a telescope
$\Gamma \vdash pat : A \Rightarrow \Delta$	declarePat	Create telescope containing all of the variables from the pattern
$\vdash a \sim b \Rightarrow \Delta$	unify	Compare two terms to create a list of definitions
$\Delta[a/x]$	doSubst	Substitute through a telescope
$\Delta_1[\bar{a}/\Delta_2]$	substTele	Substitute a list of args for the variables declared in a telescope

Figure 8: Functions for checking datatypes and case expressions

```
-- | A telescope is a list of type declarations and definitions
newtype Telescope = Telescope [Decl]
```

For example, a declaration for the `Bool` type would be the following.

```
boolDecl :: Decl
boolDecl = Data "Bool" [ConstructorDef "False" [],
                        ConstructorDef "True"  [] ]
```

9.3 Checking (simple) data constructor applications

When we have a datatype declaration, that means that new data type T of type **Type** will be added to the context. Furthermore, the context should record all of the type constructors for that type, K_i , as well as the telescope, written Δ_i , for that data constructor. This information will be used to check terms that are the applications of data constructors. For simplicity, we'll assume that data constructors must be fully applied to all of their arguments.

So our typing rule looks a little like this. We have \bar{a} as representing the list of arguments for the data constructor K .

$$\frac{\text{I-DCON-SIMPLE} \quad \begin{array}{l} K : \Delta \rightarrow T \in \Gamma \\ \Gamma \vdash \bar{a} \Leftarrow \Delta \end{array}}{\Gamma \vdash K \bar{a} \Rightarrow T}$$

We need to check that list against the telescope for the constructor, using the judgment $\Gamma \vdash \bar{a} \Leftarrow \Delta$.

In this judgement, each argument must have the right type. Furthermore, because of dependency, we substitute that argument for the variable in the rest of the telescope.

$$\frac{\text{TELE-SIG} \quad \begin{array}{l} \Gamma \vdash a : A \quad \Gamma \vdash \bar{a} \Leftarrow \Delta[a/x] \end{array}}{\Gamma \vdash a \bar{a} \Leftarrow x : A, \Delta}$$

Furthermore, we also substitute when the telescope contains definitions.

$$\frac{\text{TELE-DEF} \quad \Gamma \vdash \bar{a} \Leftarrow \Delta[a/x]}{\Gamma \vdash \bar{a} \Leftarrow x = a, \Delta}$$

When we get to the end of the argument list (i.e. there are no more arguments) we should also get to the end of the telescope.

$$\frac{\text{TELE-NIL}}{\Gamma \vdash \Leftarrow}$$

In `TypeCheck.hs`, the function `tcArgTele` essentially implements this judgement. (For reasons that we explain below, we have a special type `Arg` for the arguments to the data constructor.)

```
tcArgTele :: [Arg] -> Telescope -> TcMonad [Arg]
```

This function relies on the substitution function for telescopes, written $\Delta[a/x]$ in the rules above:

```
doSubst :: [(TName, Term)] -> [Assn] -> TcMonad [Assn]
```

9.4 Eliminating dirt simple datatypes

In your homework assignment, we used a special `if` to eliminate boolean types. Now, we'd like to replace that with the more general `case` expression, of the form `case a of { $\overline{pat_i \rightarrow a_i^i}$ }` that works with any form of datatype. What should the typing rule for that sort of expression look like?

$$\frac{\text{C-CASE-SIMPLE} \quad \begin{array}{c} \Gamma \vdash a \Rightarrow A \\ \text{whnf } A \rightsquigarrow T \quad \overline{\Gamma \vdash pat_i : T \Rightarrow \Delta_i^i} \quad \overline{\vdash a \sim pat_i \Rightarrow \Delta_i'^i} \\ \overline{\Gamma, \Delta_i, \Delta_i' \vdash a_i : A^i} \quad \text{branches exhaustive} \end{array}}{\Gamma \vdash \text{case } a \text{ of } \{ \overline{pat_i \rightarrow a_i^i} \} \Leftarrow A}$$

Mathematically, this rule checks that the scrutinee has the type of some datatype T . Then, for each case of the pattern match, this rule looks at the pattern and its type to calculate a telescope containing declarations of all of the variables bound in that pattern. That telescope is added to the context to type check each branch. Furthermore, when checking the type of each branch, we can refine that type because we know that the scrutinee is equal to the pattern.

How do we implement this rule in our language? The general strategy is as follows:

1. Infer type of the scrutinee a (`inferType`)
2. Make sure that the inferred type is some type constructor T (`ensureTCon`)
3. For each case alternative $p_i \rightarrow a_i$:

- Create the declarations Δ_i for the variables in the pattern (`declarePat`)
 - Create a list of definitions Δ'_i that follow from unifying the scrutinee a with the pattern p_i (`unify`)
 - Check the body of the case a_i in the extended context against the expected type A (`checkType`)
4. Make sure that the patterns in the cases are exhaustive (`exhaustivityCheck`)

9.5 Datatypes with parameters

The first extension of the above scheme is for *parameterized datatypes*. For example, in pi-forall we can define the `Maybe` type with the following declaration. The type parameter for this datatype `A` can be referred to in any of the telescopes for the data constructors.

```
data Maybe (A : Type) : Type where
  Nothing
  Just of (A)
```

Because this is a dependently-typed language, the variables in the telescope can be referred to later in the telescope. For example, with parameters, we can implement Sigma types as a datatype, instead of making them primitive:

```
data Sigma (A: Type) (B : A → Type) : Type
  Prod of (x:A) (B)
```

The general form of datatype declaration with parameters includes a telescope for the type constructor, as well as a telescope for each of the data constructors.

```
data T D : Type where
  Ki of Di
```

That means that when we check an occurrence of a type constructor, we need to make sure that its actual arguments match up the parameters in the telescope. For this, we can use the argument checking judgement above.

$$\frac{\text{I-TCON} \quad T : \Delta \rightarrow \mathbf{Type} \in \Gamma \quad \Gamma \vdash \bar{a} \Leftarrow \Delta}{\Gamma \vdash T \bar{a} \Rightarrow \mathbf{Type}}$$

We modify the typing rule for data constructors by marking the telescope for type constructor in the typing rule, and then substituting the actual arguments from the expected type:

$$\frac{\text{C-DCON} \quad K : \Delta_1 \rightarrow \Delta_2 \rightarrow T \in \Gamma \quad \Gamma \vdash \bar{a} \Leftarrow \Delta_2[\bar{b}/\Delta_1]}{\Gamma \vdash K \bar{a} \Leftarrow T \bar{b}}$$

For example, if we are trying to check the expression `Just True`, with expected type `Maybe Bool`, we'll first see that `Maybe` requires the telescope `(A :`

Type). That means we need to substitute `Bool` for `A` in `(_ : A)`, the telescope for `Just`. That produces the telescope `(_ : Bool)`, which we'll use to check the argument `True`.

In `TypeCheck.hs`, the function

```
substTele :: Telescope -> [ Term ] -> Telescope -> TcMonad Telescope
```

implements this operation of substituting the actual data type arguments for the parameters.

Note that by checking the type of data constructor applications (instead of inferring them) we don't need to explicitly provide the parameters to the data constructor. The type system can figure them out from the provided type.

Also note that checking mode also enables *data constructor overloading*. In other words, we can have multiple datatypes that use the same data constructor. Having the type available allows us to disambiguate.

For added flexibility we can also add code to *infer* the types of data constructors when they are not actually parameterized (and when there is no ambiguity due to overloading).

9.6 Datatypes with indices

The final step is to index our datatypes with constraints on the parameters. Indexed types let us express inductively defined relations, such as `beautiful` from Software Foundations.

```
Inductive beautiful : nat -> Prop :=
  b_0 : beautiful 0
| b_3 : beautiful 3
| b_5 : beautiful 5
| b_sum : forall n m, beautiful n -> beautiful m -> beautiful (n+m).
```

Even though `beautiful` has type `nat -> Prop`, we call `nat` this argument an index instead of a parameter because it is determined by each data constructor. It is not used uniformly in each case.

In `pi-forall`, we'll implement indices by explicitly *constraining* parameters. These constraints will just be expressed as equalities written in square brackets. In other words, we'll define `beautiful` this way:

```
data Beautiful (n : Nat) : Type where
  B0 of [n = 0]
  B3 of [n = 3]
  B5 of [n = 5]
  Bsum of (m1:Nat)(m2:Nat)(Beautiful m1)(Beautiful m2)[m = m1+m2]
```

Constraints can appear anywhere in the telescope of a data constructor. However, they are not arbitrary equality constraints—we want to consider them as deferred substitutions. So therefore, the term on the left must always be a variable.

These constraints interact with the type checker in a few places:

- When we use data constructors we need to be sure that the constraints are satisfied, by appealing to definitional equality when we are checking arguments against a telescope (in `tcArgTele`).

$$\frac{\text{TELE-DEF} \quad \Gamma \vdash \bar{a} \Leftarrow \Delta[a/x]}{\Gamma \vdash \bar{a} \Leftarrow x = a, \Delta}$$

- When we substitute through telescopes (in `doSubst`), we may need to rewrite a constraint `x = b` if we substitute for `x`.
- When we add the pattern variables to the context in each alternative of a case expression, we need to also add the constraints as definitions (see `declarePats`).

For example, if we check an occurrence of B3, i.e.

```
threeIsBeautiful : Beautiful 3
threeIsBeautiful = B3
```

this requires substituting 3 for `n` in the telescope `[n = 3]`. That produces an empty telescope.

9.6.1 Homework: Parameterized datatypes and proofs: logic

Translate the definitions and proofs in Logic chapter of Software Foundations to `pi-forall`. See `Logic.pi` for a start.

9.6.2 Homework: Indexed datatypes: finite numbers in `FinHw.pi`

The module `FinHw.pi` declares the type of numbers that are drawn from some bounded set. For example, the type `Fin 1` only includes 1 number (called Zero), `Fin 2` includes 2 numbers, etc. More generally, `Fin n` is the type of all natural numbers smaller than `n`, i.e. of all valid indices for lists of size `n`.

```
data Fin (n : Nat) : Type where
  Zero of (m:Nat)[n = Succ m]
  Succ of (m:Nat)[n = Succ m] (Fin m)
```

The file `FinHw.pi` includes a number of definitions that use these types. Two of these definitions are marked with `TRUSTME`. Replace these expressions with appropriate definitions.

9.7 Erasure and datatypes

What about putting it in data structures? We should be able to define datatypes with “specificational arguments”. For example, see `Vec.pi`.

Note: we can only erase *data* constructor arguments, not types that appear as arguments to *type* constructors. (Parameters to type constructors must always be relevant, they determine the actual type.) On the other hand, datatype

parameters are never relevant to data constructors—we don’t even represent them in the abstract syntax.

9.7.1 Homework: Erasure and Indexed datatypes: finite numbers in `FinHw.pi`

Now take your code in `FinHw.pi` and see if you can mark some of the components of the `Fin` datatype as erasable.

10 Where to go next?

10.1 Other tutorials on the implementation of dependent type systems

- A. Löb, C. McBride, W. Swierstra, *A tutorial implementation of a dependently typed lambda calculus* [LMS10].

This tutorial implements the “LambdaPi” language, and is accompanied by a Haskell implementation available online ¹⁴. The tutorial starts with an implementation of the simply-typed lambda calculus, using a locally nameless representation and bidirectional type checking, and then extends to core type system with type-in-type and uses normalization-by-evaluation for equivalence checking. It also adds natural numbers and vectors along with their eliminators, but does not include a general form of (inductive) datatypes.

- Lennart Augustsson. *Simple, Easier!* [Aug07]

The goal of Augustsson’s tutorial, inspired by an earlier version of LambdaPi, is to be the simplest implementation possible. Therefore, he uses a string representation of variables and binding and weak-head normalization for implementation of equality. His version is a Haskell implementation of Barendregt’s lambda cube [Bar93].

- Coquand, Kinoshita, Nordstrom, Takeyama. *A simple type-theoretic language: Mini-TT* [CKNT09]

Coquand et al. describe another Haskell implementation of dependent types, called “Mini-TT”, that shows how to implement (in about 400 lines of code) a small dependently-typed language that includes Σ -types, void, unit and sums. The language does not include propositional equality or indexed datatypes, nor does it enforce termination. The implementation uses strings to represent variable binding, which are converted to de Bruijn indices during normalization-by-evaluation.

- Andrej Bauer, *How to implement dependent type theory* [Bau12]

Bauer explores several different implementations of dependent type theory

¹⁴<https://www.andres-loeh.de/LambdaPi/>

in a series of blog posts using the OCaml language. The first version uses a named representation (i.e. strings) and generates fresh names during substitution and alpha-equality. This version uses full normalization to implement definitional equality. Then, in the second version, Bauer revises the implementation to use normalization by evaluation. In the third version, he revises again to use de Bruijn indices, explicit substitutions and switches back to weak-head normalization.

- Tiark Rompf, *Implementing Dependent Types* [Rom20].

A recent blog post tutorial, Rompf uses Javascript to implement a core dependent type theory using higher-order abstract syntax.

- Andras Kovacs, *Elaboration Zoo* [Kov]

Kovacs' github repository includes several different implementations of evaluation and type checking for a minimal dependently typed languages. These implementations are accompanied by video recordings of a seminar that goes over this code. Notably, this tutorial includes implicit argument synthesis via (higher-order) unification.

10.2 Related work for `pi-forall`

Section 3: Core Dependent Types Because they lack a rule for conversion, the typing rules for the core dependent type system presented in Section 3 do not correspond to any existing type system. However, jumping ahead, the core language with the inclusion of conversion (rule T-CONV) has been well-studied.

As a foundation of logic, the inclusion of rule rule T-TYPE makes the system inconsistent. Indeed, Martin-Löf's original formulation of type theory [ML71] included this rule until the resulting paradox was demonstrated by Girard. This inconsistency was resolved in later versions of the systems [ML75] and in the Calculus of Constructions [?] by stratifying **Type**: first into two levels and then into an infinite hierarchy of universes [Luo90].

The inconsistency of this system as a logic does not prevent its use as a programming language. Cardelli [Car86] showed that this language was type sound and explored applications for programming.

Section 4: Bidirectional type checking The expressiveness of dependently-typed programming languages put them out of reach of the Hindley Milner type inference algorithm found in Haskell and ML. As a result, type checkers and elaborators for these languages must abandon complete type inference and require users to annotate their code.

Bidirectional type systems provide a platform for propagating type annotations through judgements, and notably, support a syntax-directed specification of where to apply the conversion rule during type checking. As a result, such type systems are frequently used as the basis for type checking, with elaboration for implicit arguments layered on top.

Pierce and Turner [PT00] popularized the use of bidirectional type systems, in the context of polymorphic languages with subtyping. The Haskell language uses bidirectional type checking layered on top of Hindley-Milner unification in order to support type inference for higher-rank polymorphism [PVWS07]. This type system extension allows functions to take polymorphic functions as arguments. Notably, the paper includes a well-documented reference implementation of the type system in the appendix.

David Christiansen’s tutorial [Chr13] introduces the idea of bidirectional type checking as an independent topic. For another description of the role of bidirectional type systems in the context of dependent types, see Löh, McBride and Swierstra’s tutorial [LMS10]. Unlike `pi-forall`, this tutorial separates the syntax of terms based whether they are inferable or checkable.

Finally, Dunfield and Krishnaswami’s extensive survey paper [DK21] catalogs and organizes recent results in this area.

Section 6: Definitional equality Equality has been called the hardest problem in type theory, and that is not an exaggeration.

Many implementations, such as Coq and Agda use a procedure called *normalization-by-evaluation* to decide whether two terms are equivalent.

- Coq [GL02]
- Andreas Abel’s habilitation thesis [Abe13].
- Daniel Gratzer, NBE for MLTT <https://github.com/jozefg/nbe-for-mltt>

Section 8: Compile-time and run-time irrelevance The implementation of irrelevance in `pi-forall` is inspired by the Dependent Dependency Calculus (DDC) [CEIW22], a pure type system parameterized by a dependency lattice. The selection of the lattice determines the precise form of irrelevance described by the type system: does the system include run-time irrelevance only, compile-time irrelevance only, both sorts of irrelevance combined together (as in `pi-forall`), or both sorts and allow them to be distinguished.

Compile-time irrelevance only Pfenning [Pfe01] used a modal type system to identify arguments that can be ignored when checking equivalence. This system, extended by Abel and Scherer [AS12], became the basis for Agda’s implementation of compile-time irrelevance.

Run-time irrelevance only The Coq proof assistance distinguishes the `Prop` universe from the `Type` universe for several reasons. However, one feature of this distinction is that all terms in the `Prop` universe can be erased prior to execution, leading to faster evaluation of extracted code.

More recently, *quantitative type theory* [Atk18] unifies run-time irrelevance and linear type systems by restricting the number of times that variables may

be used in a given context. This resulting system has been adopted for the implementations of irrelevance in the Agda and Idris [Bra21] languages.

Both run-time and compile-time irrelevance, but no distinction The Implicit Calculus of Constructions [Miq01] ensures that some arguments are treated irrelevantly (at both runtime and compiletime) by simply not including them in the syntax of terms. This work was extended to decidable type checking through the use of an erasure operation—terms may include irrelevant arguments, but these arguments must be erased prior to runtime and equality checking [BB08]. This approach led to Erasure Pure Type Systems (EPTS) [MLS08], and the treatment of irrelevance in Dependent Haskell [WVdAE17].

Section 7: Dependently-typed pattern matching The implementation of indexed datatypes in `pi-forall` is inspired by the compilation of *Generalized Algebraic Datatypes* in Haskell [VPJS11, VWPJ06]. The distinguishing feature of this formalization is that internally, datatype indices are represented by constrained parameters and so are an extension of the parameterized recursive sums of products found in functional programming languages.

In contrast, the theory of inductive datatypes found in Coq and Agda must ensure that they cannot be used to implement nonterminating functions. This means that there are several restrictions: they must be strictly positive, their eliminators must be defined separately, and they must satisfy constraints based on universe-levels.

- Coq pattern matching: Coq User manual
- Agda pattern matching: Ulf Norell’s dissertation

11 Conclusion

References

- [Abel13] Andreas Abel. *Normalization by Evaluation: Dependent Types and Impredicativity*. Habilitation thesis, Ludwig-Maximilians-Universität München, 2013.
- [AS12] Andreas Abel and Gabriel Scherer. On irrelevance and algorithmic equality in predicative type theory. *Logical Methods in Computer Science*, 8(1), mar 2012.
- [Atk18] Robert Atkey. Syntax and semantics of quantitative type theory. In *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science*, LICS ’18, page 56–65, New York, NY, USA, 2018. Association for Computing Machinery.
- [Aug07] Lennart Augustsson. `Simpler, easier!`, 2007. Avail-

- able from <https://augustss.blogspot.com/2007/10/simpler-easier-in-recent-paper-simply.html>.
- [Bar93] H. P. Barendregt. *Lambda Calculi with Types*, page 117–309. Oxford University Press, Inc., USA, 1993.
 - [Bau12] Andrej Bauer. How to implement dependent type theory, November 2012. Available from <http://math.andrej.com/2012/11/08/how-to-implement-dependent-type-theory-i/>.
 - [BB08] Bruno Barras and Bruno Bernardo. The implicit calculus of constructions as a programming language with dependent types. In Roberto Amadio, editor, *Foundations of Software Science and Computational Structures*, pages 365–379, Berlin, Heidelberg, 2008. Springer Berlin Heidelberg.
 - [Bra21] Edwin C. Brady. Idris 2: Quantitative type theory in practice. In Anders Møller and Manu Sridharan, editors, *35th European Conference on Object-Oriented Programming, ECOOP 2021, July 11-17, 2021, Aarhus, Denmark (Virtual Conference)*, volume 194 of *LIPICs*, pages 9:1–9:26. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2021.
 - [Car86] Luca Cardelli. A polymorphic λ calculus with Type:Type. Technical Report 10, Digital Equipment Corporation, SRC, 1986.
 - [CEIW22] Pritam Choudhury, Harley Eades III, and Stephanie Weirich. A dependent dependency calculus. In Ilya Sergey, editor, *Programming Languages and Systems, ESOP 2022*, volume 13240 of *Lecture Notes in Computer Science*, pages 403–430, Cham, 2022. Springer International Publishing.
 - [Chr13] David Raymond Christiansen. Bidirectional typing rules: A tutorial. Tutorial, October 2013.
 - [CKNT09] Thierry Coquand, Yoshiki Kinoshita, Bengt Nordström, and Makoto Takeyama. A simple type-theoretic language: Mini-tt. In Yves Bertot, Gérard Huet, Jean-Jacques Lévy, and Gordon Plotkin, editors, *From Semantics to Computer Science: Essays in Honour of Gilles Kahn*, page 139–164. Cambridge University Press, 2009.
 - [dB72] Nicolaas Govert de Bruijn. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the Church-Rosser theorem. *Indagationes Mathematicae (Proceedings)*, 75(5):381–392, 1972.
 - [DK21] Jana Dunfield and Neel Krishnaswami. Bidirectional typing. *ACM Comput. Surv.*, 54(5), may 2021.

- [GL02] Benjamin Grégoire and Xavier Leroy. A compiled implementation of strong reduction. In *Proceedings of the Seventh ACM SIGPLAN International Conference on Functional Programming*, ICFP '02, page 235–246, New York, NY, USA, 2002. Association for Computing Machinery.
- [Har16] Robert Harper. *Practical Foundations for Programming Languages*. Cambridge University Press, USA, 2nd edition, 2016.
- [Kov] Andras Kovacs. Elaboration zoo. 2022. <https://github.com/AndrasKovacs/elaboration-zoo/>.
- [LMS10] Andres Löb, Conor McBride, and Wouter Swierstra. A tutorial implementation of a dependently typed lambda calculus. *Fundam. Inform.*, 102:177–207, 01 2010. Available from <http://www.andres-loeb.de/LambdaPi/>.
- [Luo90] Zhaohui Luo. *An Extended Calculus of Constructions*. PhD thesis, University of Edinburgh, 1990.
- [Miq01] Alexandre Miquel. The implicit calculus of constructions extending pure type systems with an intersection type binder and subtyping. In Samson Abramsky, editor, *Typed Lambda Calculi and Applications*, pages 344–359, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- [ML71] Per Martin-Löf. A theory of types. Unpublished manuscript, 1971.
- [ML75] Per Martin-Löf. An intuitionistic theory of types: predicative part. In H.E. Rose and J.C. Shepherdson, editors, *Logic Colloquium '73, Proceedings of the Logic Colloquium*, volume 80 of *Studies in Logic and the Foundations of Mathematics*, pages 73–118. North-Holland, 1975.
- [MLS08] Nathan Mishra-Linger and Tim Sheard. Erasure and polymorphism in pure type systems. In *Proceedings of the Theory and Practice of Software, 11th International Conference on Foundations of Software Science and Computational Structures*, FOSACS'08/ETAPS'08, page 350–364, Berlin, Heidelberg, 2008. Springer-Verlag.
- [Pfe01] Frank Pfenning. Intensionality, extensionality, and proof irrelevance in modal type theory. In *Proceedings of the 16th Annual IEEE Symposium on Logic in Computer Science*, LICS '01, pages 221–, Washington, DC, USA, 2001. IEEE Computer Society.
- [Pie02] Benjamin C. Pierce. *Types and Programming Languages*. The MIT Press, 1st edition, 2002.

- [PT00] Benjamin C. Pierce and David N. Turner. Local type inference. *ACM Trans. Program. Lang. Syst.*, 22(1):1–44, January 2000.
- [PVWS07] Simon Peyton Jones, Dimitrios Vytiniotis, Stephanie Weirich, and Mark Shields. Practical type inference for arbitrary-rank types. *Journal of Functional Programming*, 17(1), January 2007.
- [Rom20] Tiark Rompf. Implementing dependent types, December 2020. Available from <https://tiarkrompf.github.io/notes/?/dependent-types/>.
- [SNO⁺07] Peter Sewell, Francesco Zappa Nardelli, Scott Owens, Gilles Peskine, Thomas Ridge, Susmit Sarkar, and Rok Strniša. Ott: Effective tool support for the working semanticist. In *Proceedings of the 12th ACM SIGPLAN International Conference on Functional Programming*, ICFP ’07, page 1–12, New York, NY, USA, 2007. Association for Computing Machinery.
- [Ste17] Guy L. Steele. It’s time for a new old language. In *Proceedings of the 22nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPOPP ’17, page 1, New York, NY, USA, 2017. Association for Computing Machinery.
- [VPJSS11] Dimitrios Vytiniotis, Simon Peyton Jones, Tom Schrijvers, and Martin Sulzmann. Outsidein(x): Modular type inference with local assumptions. *Journal of Functional Programming*, 21:333–412, September 2011.
- [VWPJ06] Dimitrios Vytiniotis, Stephanie Weirich, and Simon Peyton Jones. Simple unification-based type inference for gadt. In *International Conference on Functional Programming (ICFP’06)*. ACM SIGPLAN, April 2006. 2016 ACM SIGPLAN Most Influential ICFP Paper Award.
- [WVdAE17] Stephanie Weirich, Antoine Voizard, Pedro Henrique Avezedo de Amorim, and Richard A. Eisenberg. A specification for dependent types in Haskell. *Proc. ACM Program. Lang.*, 1(ICFP):31:1–31:29, August 2017.